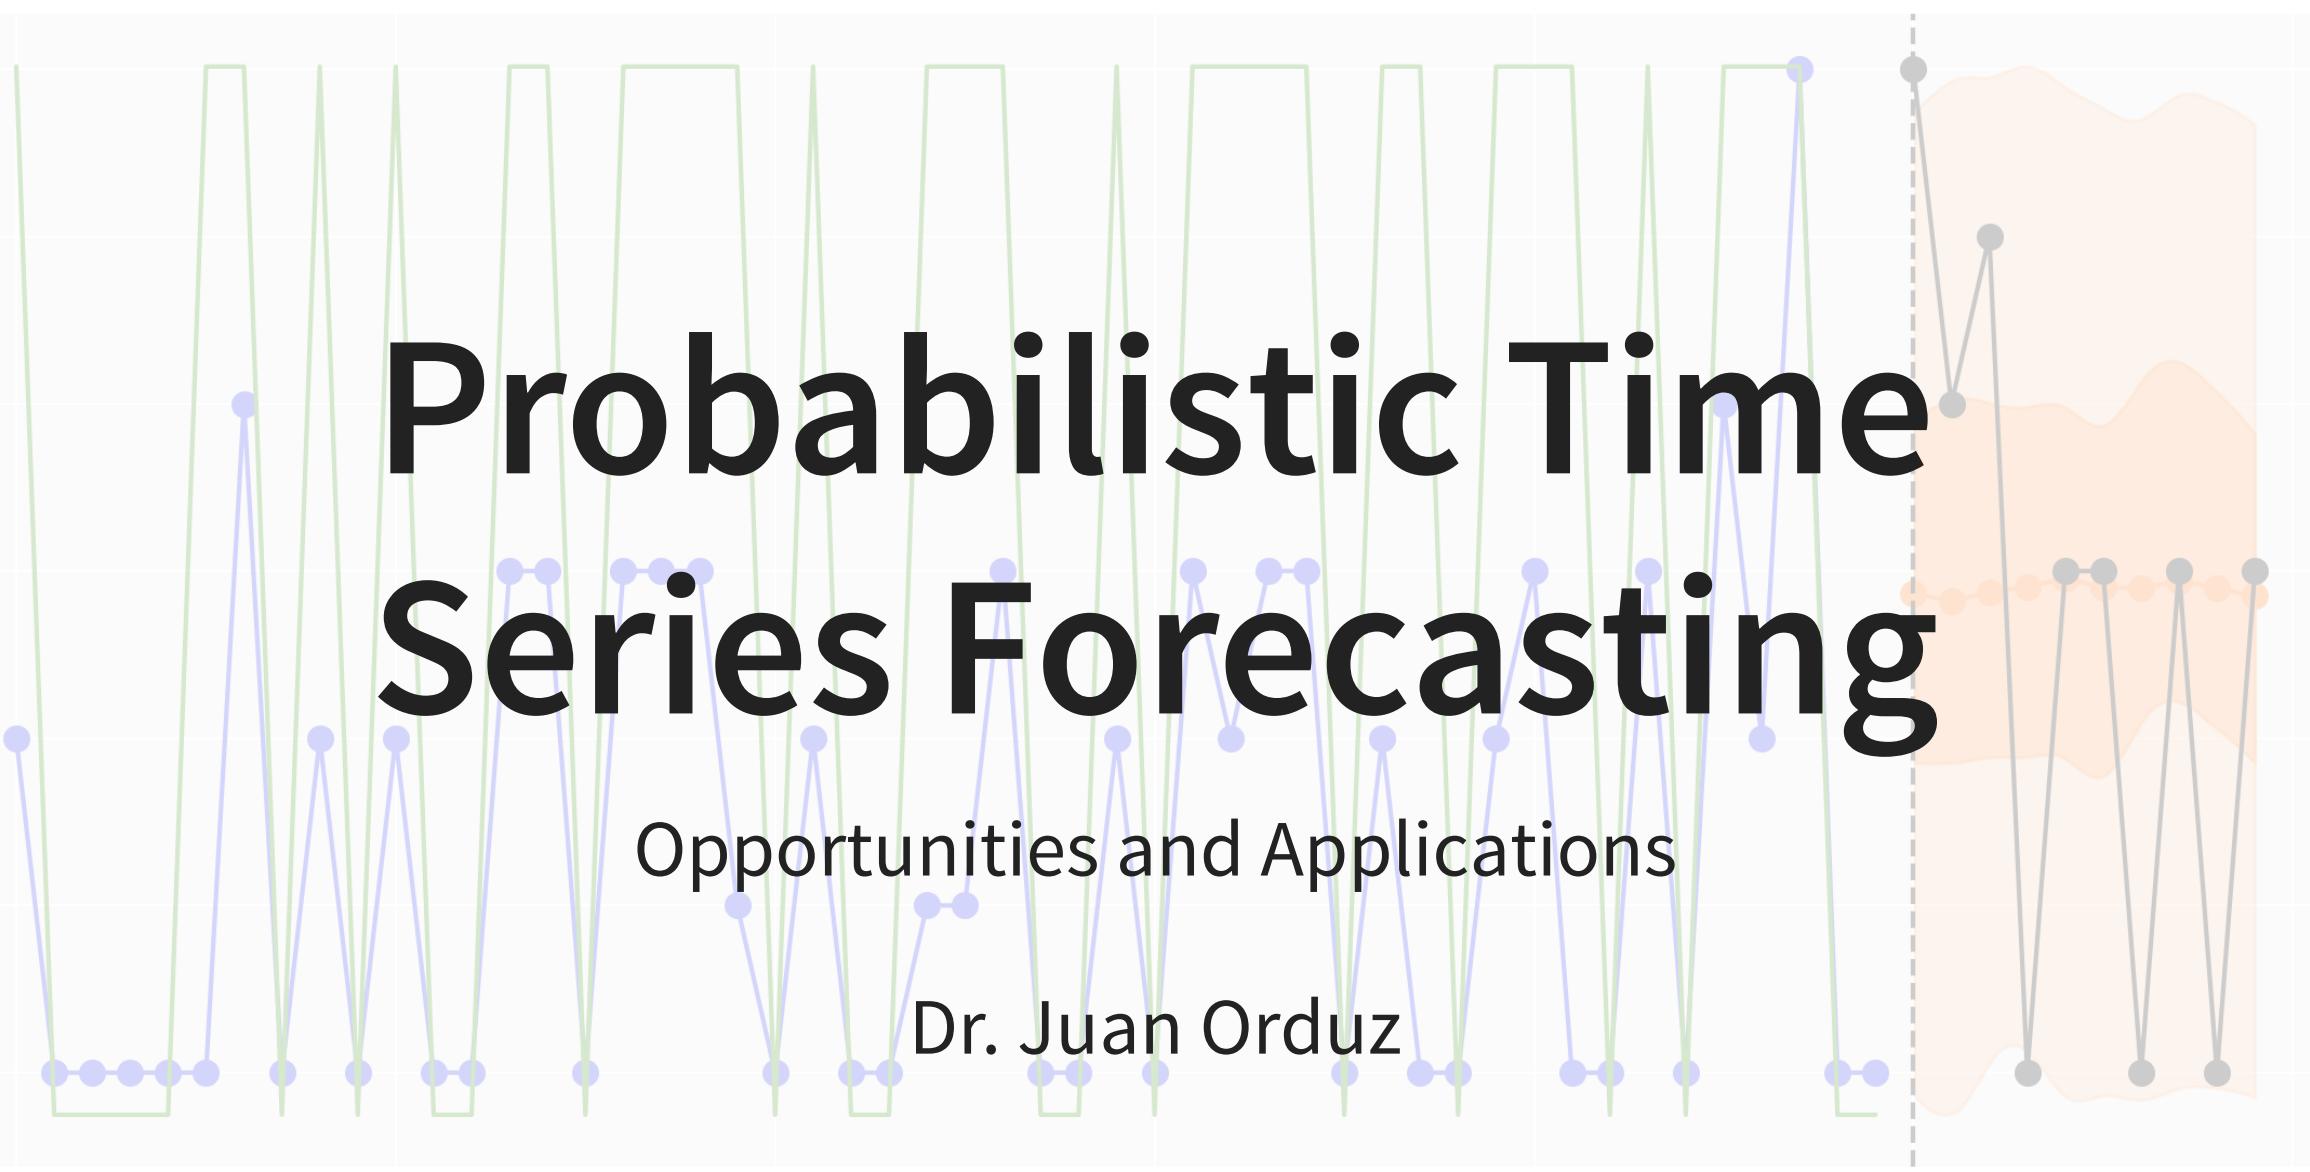


Probabilistic Time Series Forecasting

Opportunities and Applications

Dr. Juan Orduz



94% HDI
50% HDI

Posterior Predictive Mean
Test (observed)

Available Training



Outline

- 1. Introduction
- 2. Scan in NumPyro
- 3. Exponential Smoothing
- 4. ARIMA
- 5. Hierarchical Models
- 6. Intermittent Demand
- 7. Censored Demand
- 8. Price Elasticities
- 9. Prior Model Calibration
- 10. References



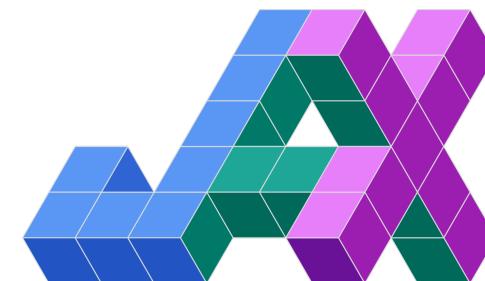
Why Probabilistic Forecasting?



- Interpretability
 - Trust on the results
- Uncertainty quantification
 - Risk Assessment
 - Decision Making
- Customization
 - Feature engineering
 - Special constraints
 - Calibrate with domain knowledge
- Scale
 - Good results on production environments
 - Take advantage of GPU



[NumPyro](#) is a lightweight probabilistic programming library that provides a NumPy backend for [Pyro](#). It relies on [JAX](#) for automatic differentiation and JIT compilation to GPU / CPU.



Pyro Forecasting Module 🔥



stable

Search docs

PYRO CORE:

- Getting Started
- Primitives
- Inference
- Distributions
- Parameters
- Neural Networks
- Optimization
- Poutine (Effect handlers)
- Miscellaneous Ops
- Settings
- Testing Utilities

CONTRIBUTED CODE:

- Automatic Name Generation
- Bayesian Neural Networks
- Causal Effect VAE
- Easy Custom Guides
- Epidemiology
- Pyro Examples

☰ Forecasting

[Home](#) » Forecasting

[Edit on GitHub](#)

Forecasting

`pyro.contrib.forecast` is a lightweight framework for experimenting with a restricted class of time series models and inference algorithms using familiar Pyro modeling syntax and PyTorch neural networks.

Models include hierarchical multivariate heavy-tailed time series of ~1000 time steps and ~1000 separate series. Inference combines subsample-compatible variational inference with Gaussian variable elimination based on the `GaussianHMM` class. Inference using Hamiltonian Monte Carlo sampling is also supported with `HMCForecaster`. Forecasts are in the form of joint posterior samples at multiple future time steps.

Hierarchical models use the familiar `plate` syntax for general hierarchical modeling in Pyro. Plates can be subsampled, enabling training of joint models over thousands of time series. Multivariate observations are handled via multivariate likelihoods like `MultivariateNormal`, `GaussianHMM`, or `LinearHMM`. Heavy tailed models are possible by using `StudentT` or `Stable` likelihoods, possibly together with `LinearHMM` and reparameterizers including `StudentTReparam`, `StableReparam`, and `LinearHMMReparam`.

Seasonality can be handled using the helpers `periodic_repeat()`, `periodic_cumsum()`, and `periodic_features()`.

See `pyro.contrib.timeseries` for ways to construct temporal Gaussian processes useful as likelihoods.

For example usage see:



Kyle Caron

Modeling Anything With First Principles: Demand under extreme stockouts

TIME SERIES

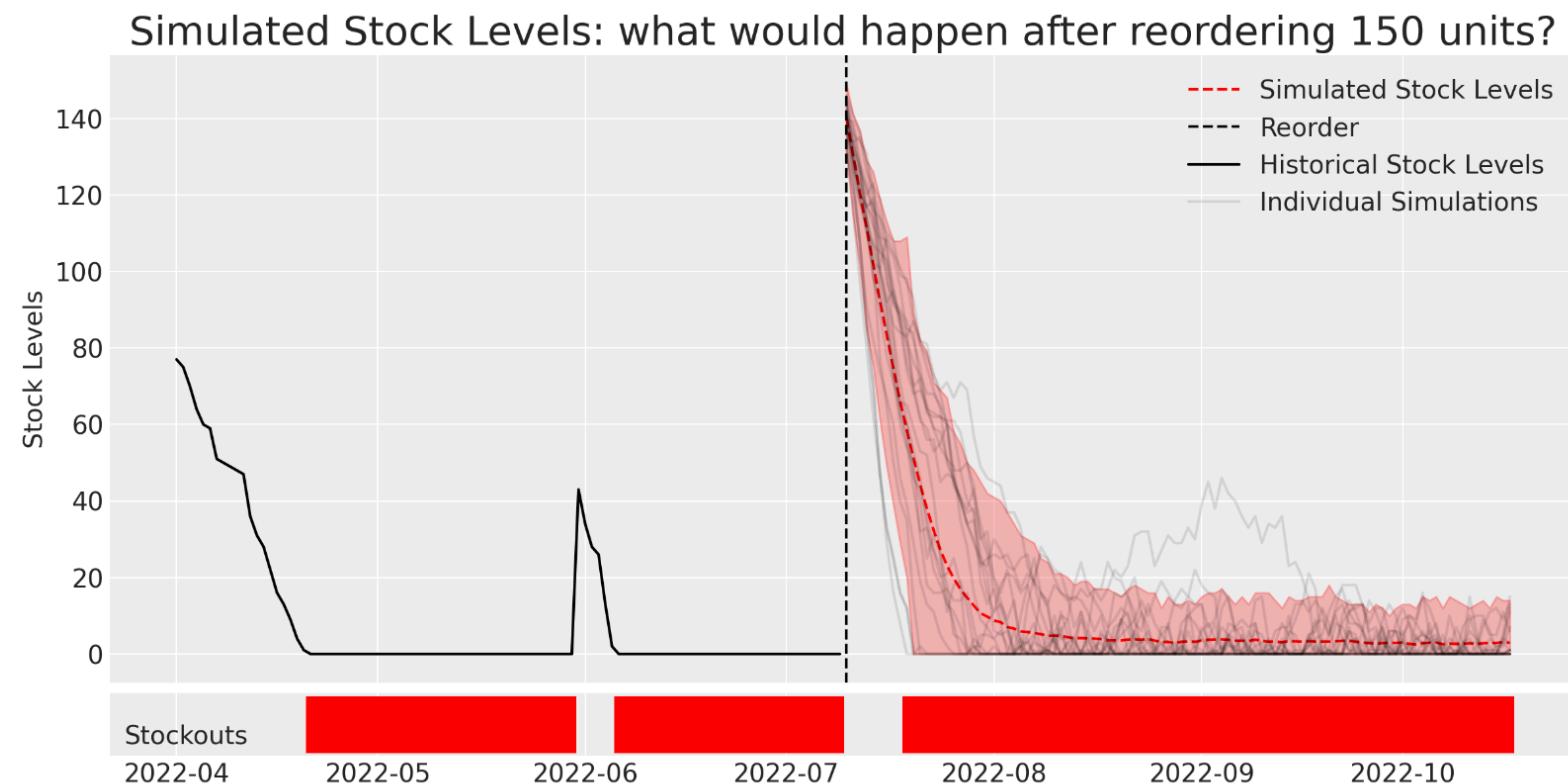
DEMAND MODELING

CAUSAL INFERENCE

SUPPLY CHAIN

DISCRETE CHOICE

SURVIVAL ANALYSIS



NumPyro - SGT Example Model 😊

Model

The model we are going to use is called **Seasonal, Global Trend**, which when tested on 3003 time series of the [M-3 competition](#), has been known to outperform other models originally participating in the competition:

$$\begin{aligned}\text{exp-val}_t &= \text{level}_{t-1} + \text{coef-trend} \times \text{level}_{t-1}^{\text{pow-trend}} + s_t \times \text{level}_{t-1}^{\text{pow-season}}, \\ \sigma_t &= \sigma \times \text{exp-val}_t^{\text{powx}} + \text{offset}, \\ y_t &\sim \text{StudentT}(\nu, \text{exp-val}_t, \sigma_t)\end{aligned}$$

, where `level` and `s` follows the following recursion rules:

$$\begin{aligned}\text{level-p} &= \begin{cases} y_t - s_t \times \text{level}_{t-1}^{\text{pow-season}} & \text{if } t \leq \text{seasonality}, \\ \text{Average}[y(t-\text{seasonality}+1), \dots, y(t)] & \text{otherwise,} \end{cases} \\ \text{level}_t &= \text{level-sm} \times \text{level-p} + (1 - \text{level-sm}) \times \text{level}_{t-1}, \\ s_{t+\text{seasonality}} &= \text{s-sm} \times \frac{y_t - \text{level}_t}{\text{level}_{t-1}^{\text{pow-trend}}} + (1 - \text{s-sm}) \times s_t.\end{aligned}$$



...



Scan ★

An efficient implementation of **for** loops

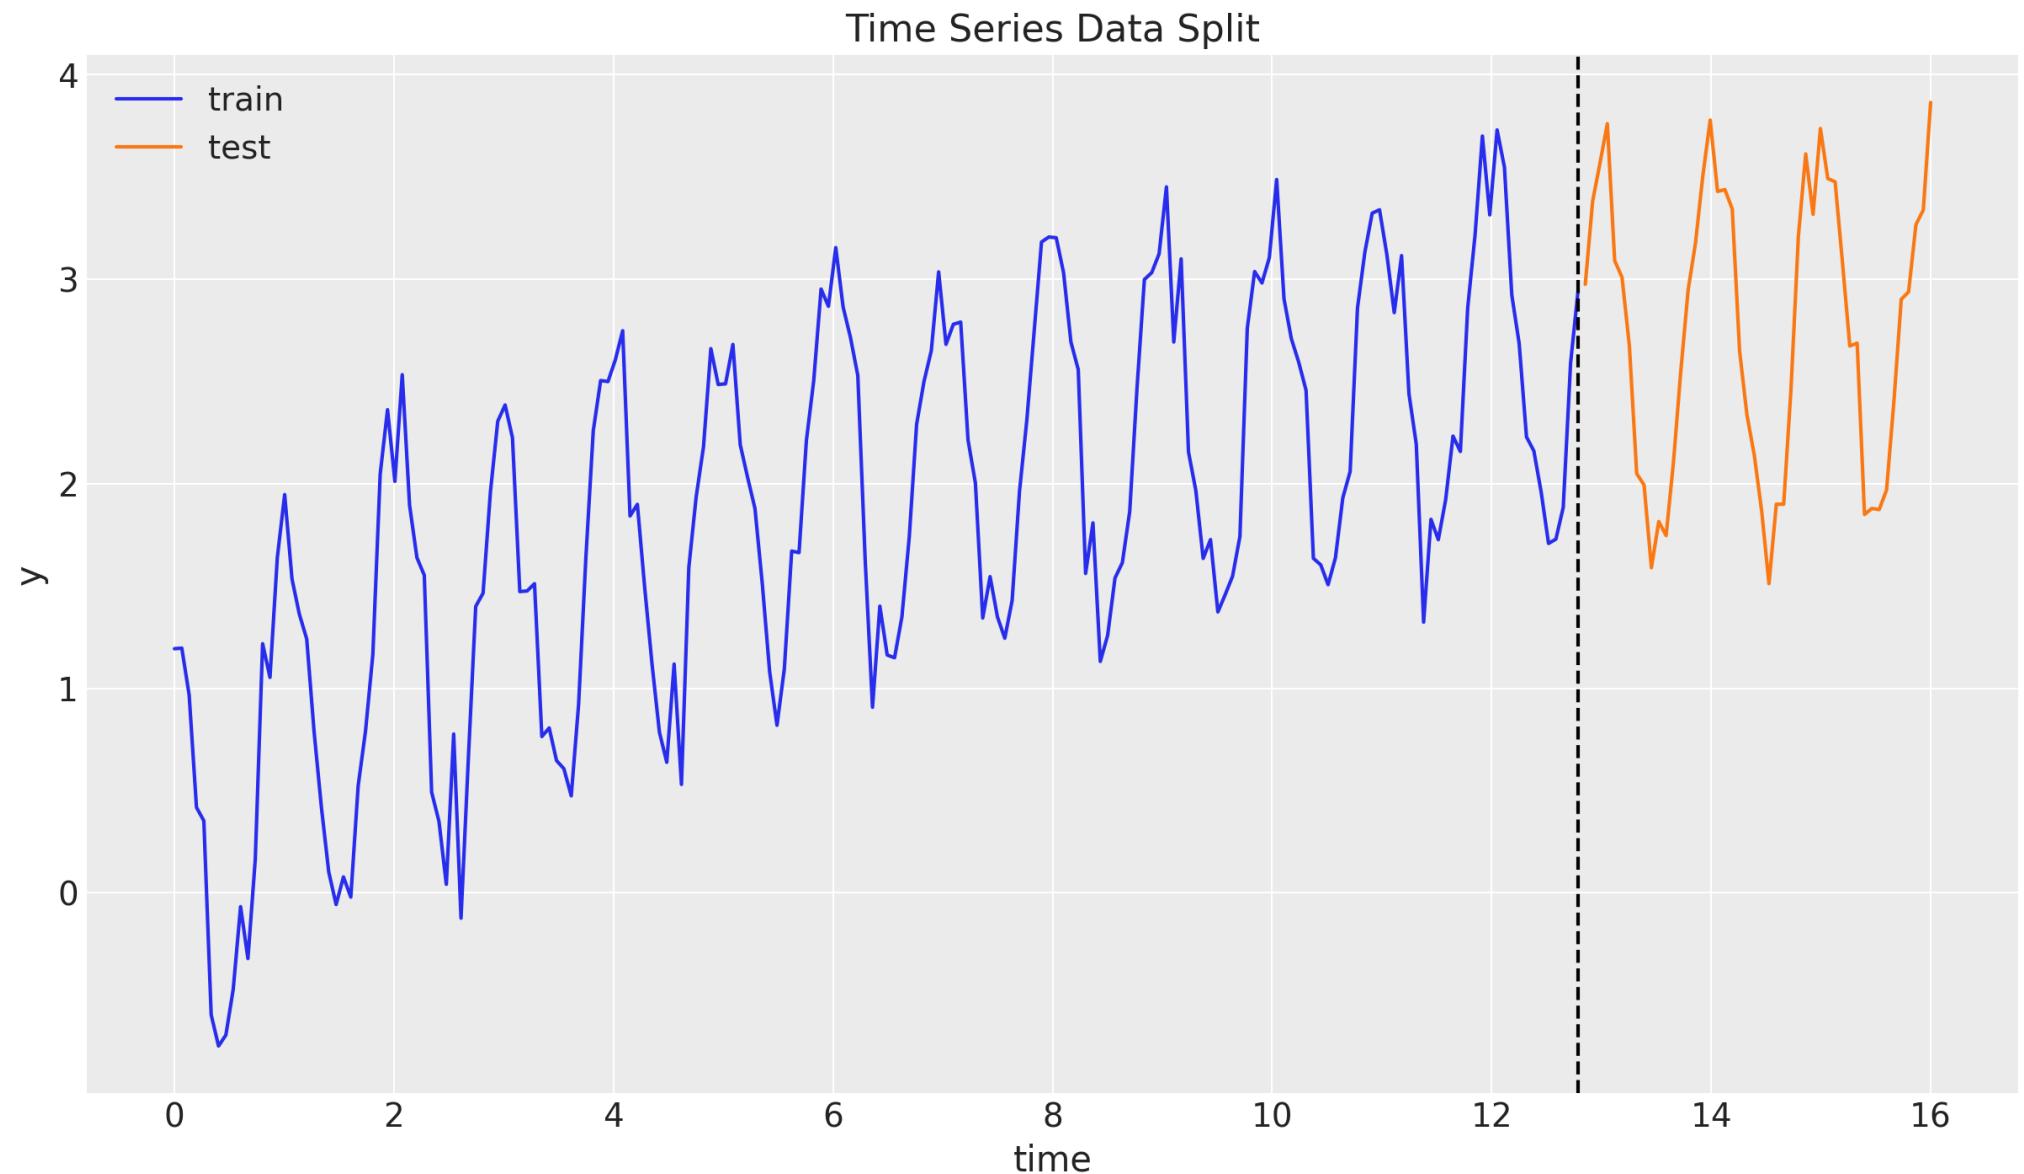
```
1 def scan(f, init, xs):
2     """Pure Python implementation of scan.
3
4     Parameters
5     -----
6     f : A a Python function to be scanned.
7     init : An initial loop carry value
8     xs : The value over which to scan along the leading axis.
9     """
10    carry = init
11    ys = []
12    for x in xs:
13        carry, y = f(carry, x)
14        ys.append(y)
15    return carry, np.stack(ys)
```



We need recursive relationships! $y_t \longmapsto y_{t+1}$



Example: Exponential Smoothing



Example: Exponential Smoothing

$$\hat{y}_{t+h|t} = l_t$$

$$l_t = \alpha y_t + (1 - \alpha)l_{t-1}$$

- y_t is the observed value at time t .
- $\hat{y}_{t+h|t}$ is the forecast of the value at time $t + h$ given the information up to time t .
- l_t is the level at time t .
- α is the smoothing parameter. It is a value between 0 and 1.



Example: Exponential Smoothing

$$\hat{y}_{t+h|t} = l_t$$

$$l_t = \alpha y_t + (1 - \alpha)l_{t-1}$$

```
1  def transition_fn(carry, t):
2      previous_level = carry
3
4      level = jnp.where(
5          t < t_max,
6          level_smoothing * y[t] + (1 - level_smoothing) * previous_level,
7          previous_level,
8      )
9
10     mu = previous_level
11     pred = numpyro.sample("pred", dist.Normal(loc=mu, scale=noise))
12
13     return level, pred
```



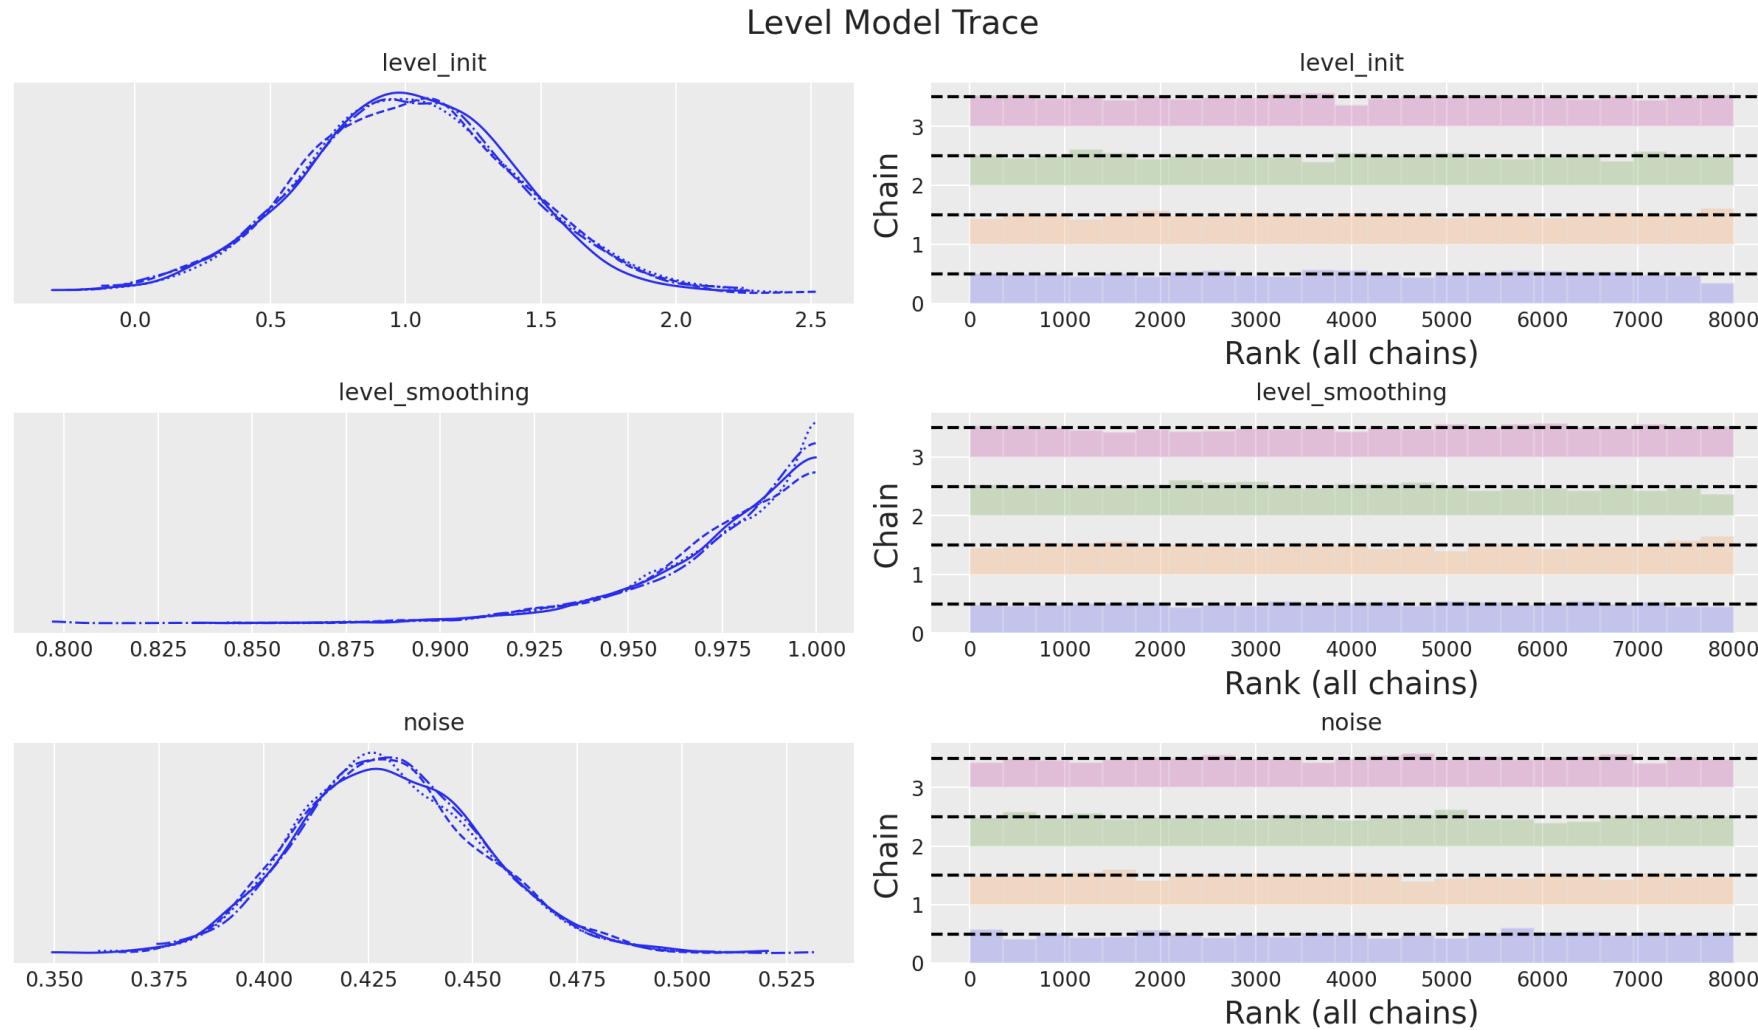
Example: Exponential Smoothing

```
1 def level_model(y: ArrayLike, future: int = 0) -> None:
2     # Get time series length
3     t_max = y.shape[0]
4
5     # --- Priors ---
6     ## Level
7     level_smoothing = numpyro.sample(
8         "level_smoothing", dist.Beta(concentration1=1, concentration0=1)
9     )
10    level_init = numpyro.sample("level_init", dist.Normal(loc=0, scale=1))
11
12    ## Noise
13    noise = numpyro.sample("noise", dist.HalfNormal(scale=1))
14
15    # --- Transition Function ---
16    def transition_fn(carry, t):
17        ...
18
19    # --- Run Scan ---
20    with numpyro.handlers.condition(data={"pred": y}):
21        _, preds = scan(
```

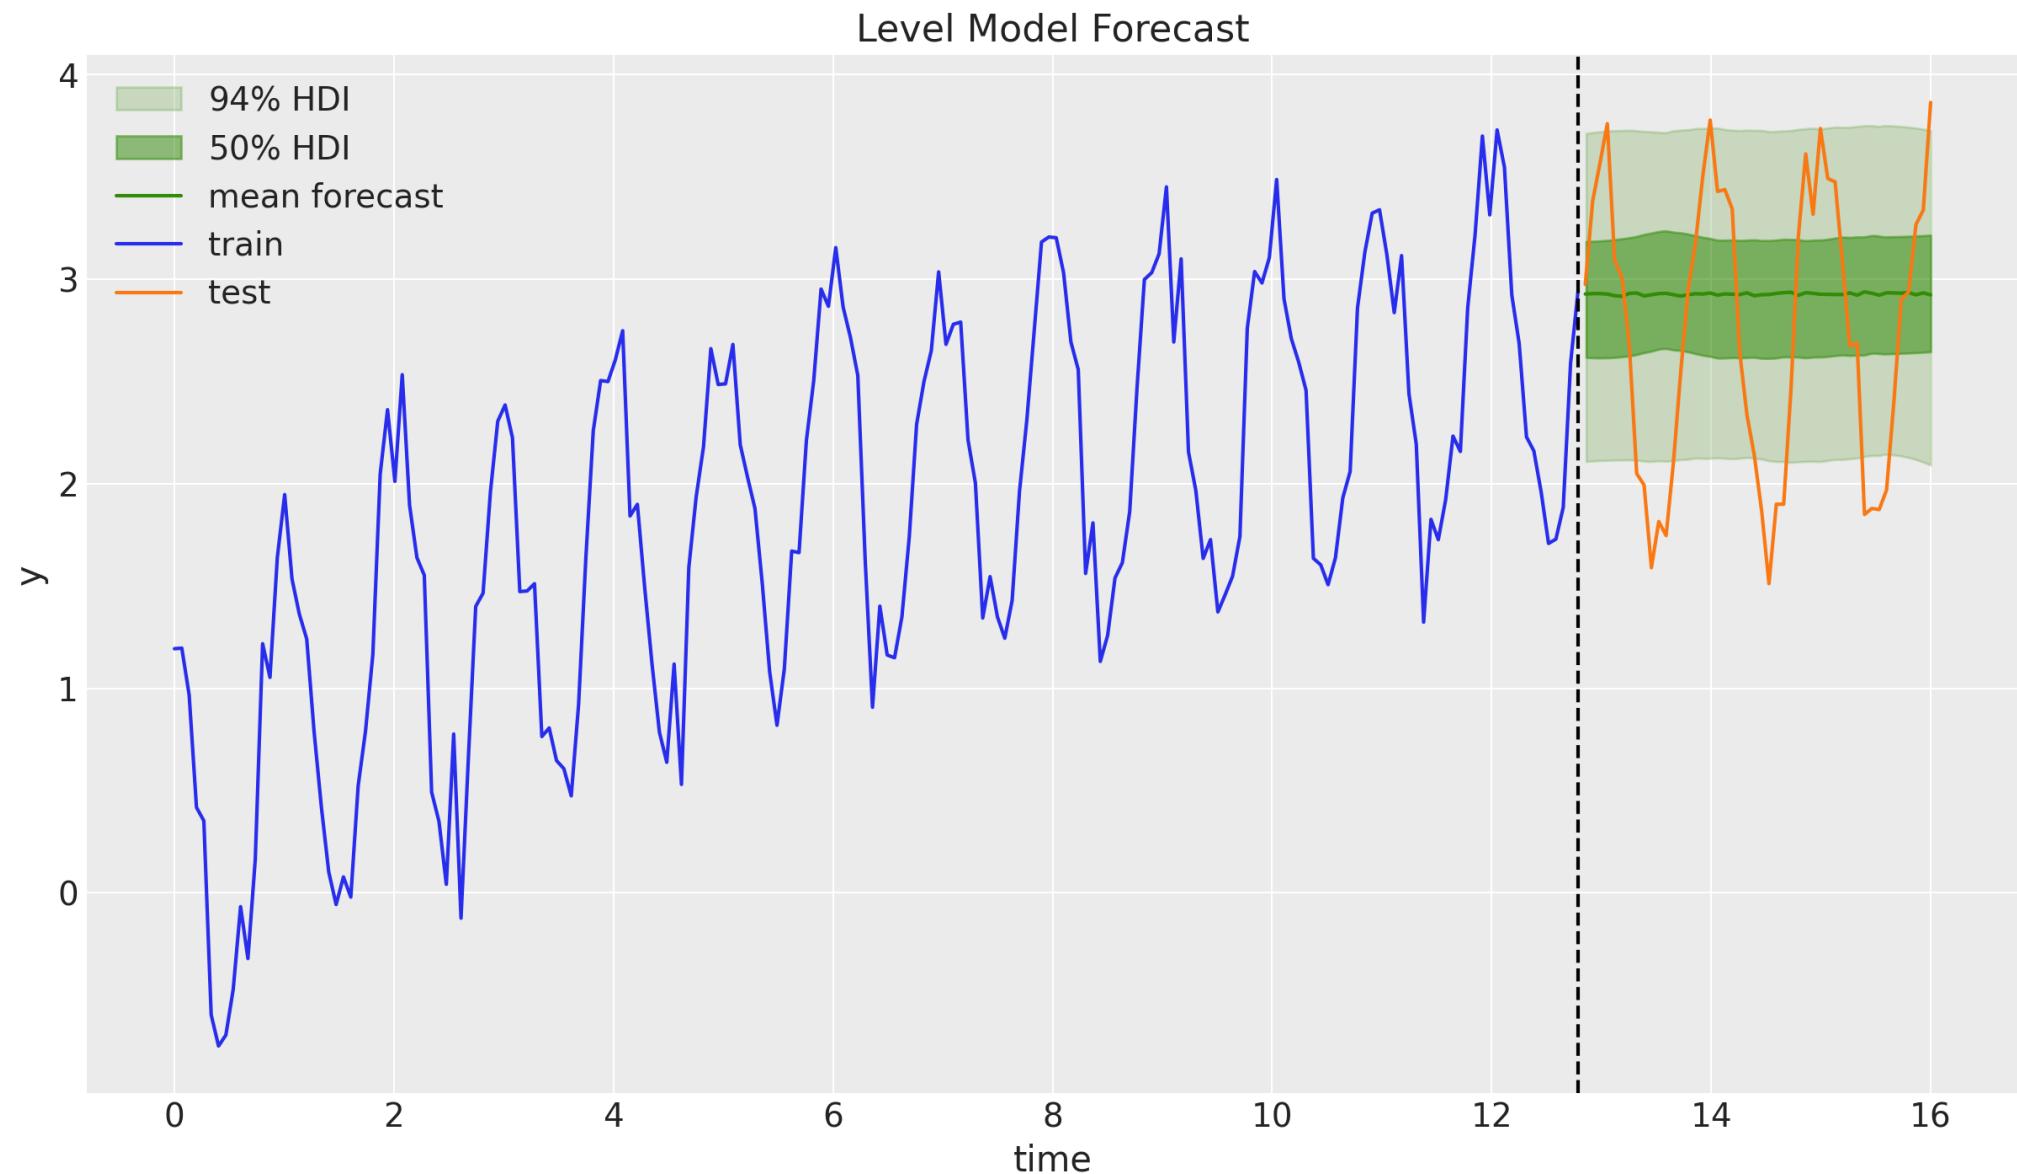


Example: Exponential Smoothing

Posterior Distribution Parameters

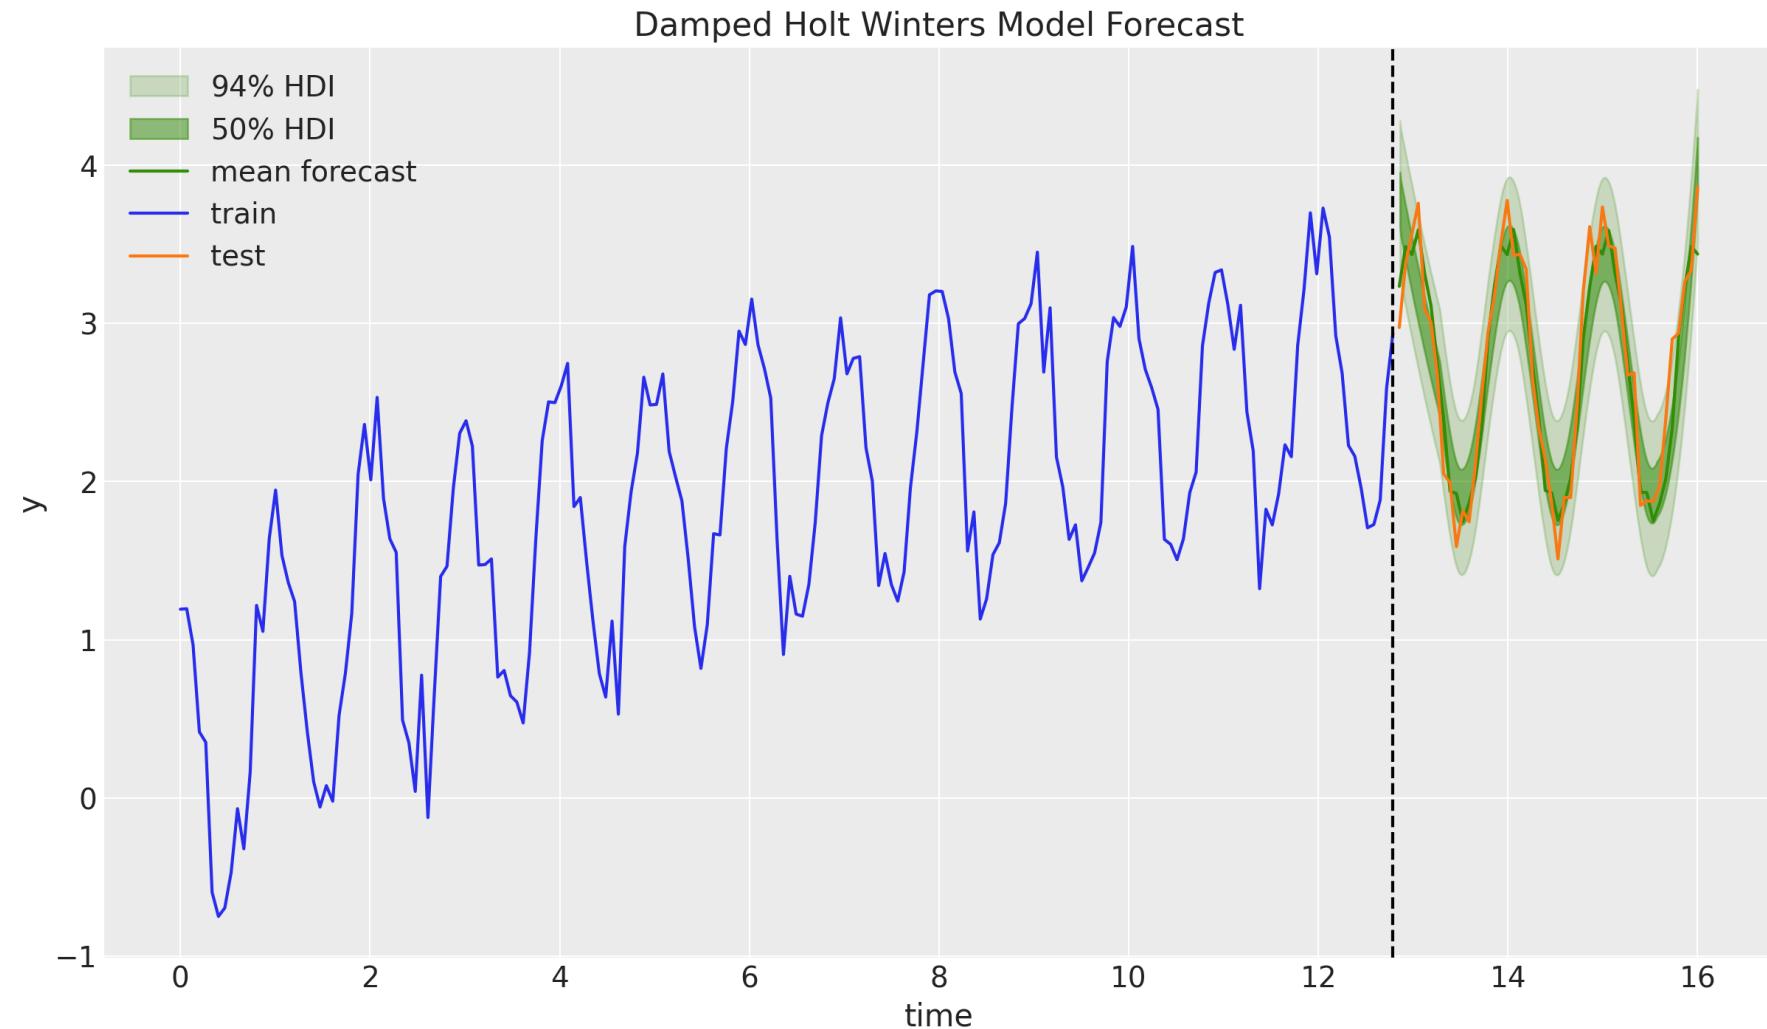


Example: Exponential Smoothing



Example: Exponential Smoothing

Trend + Seasonal + Damped

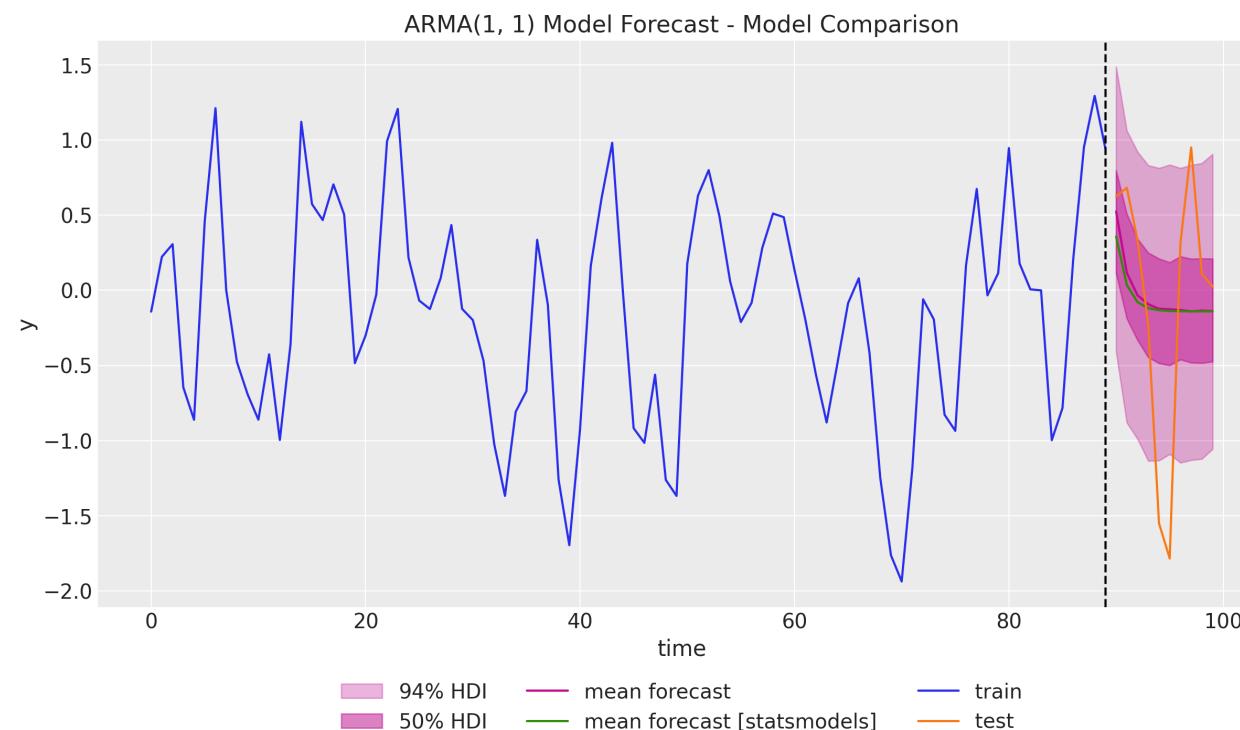


Example: ARMA(1, 1) Model

```

1  def transition_fn(carry, t):
2      y_prev, error_prev = carry
3      ar_part = phi * y_prev
4      ma_part = theta * error_prev
5      pred = mu + ar_part + ma_part
6      error = y[t] - pred
7      return (y[t], error), error

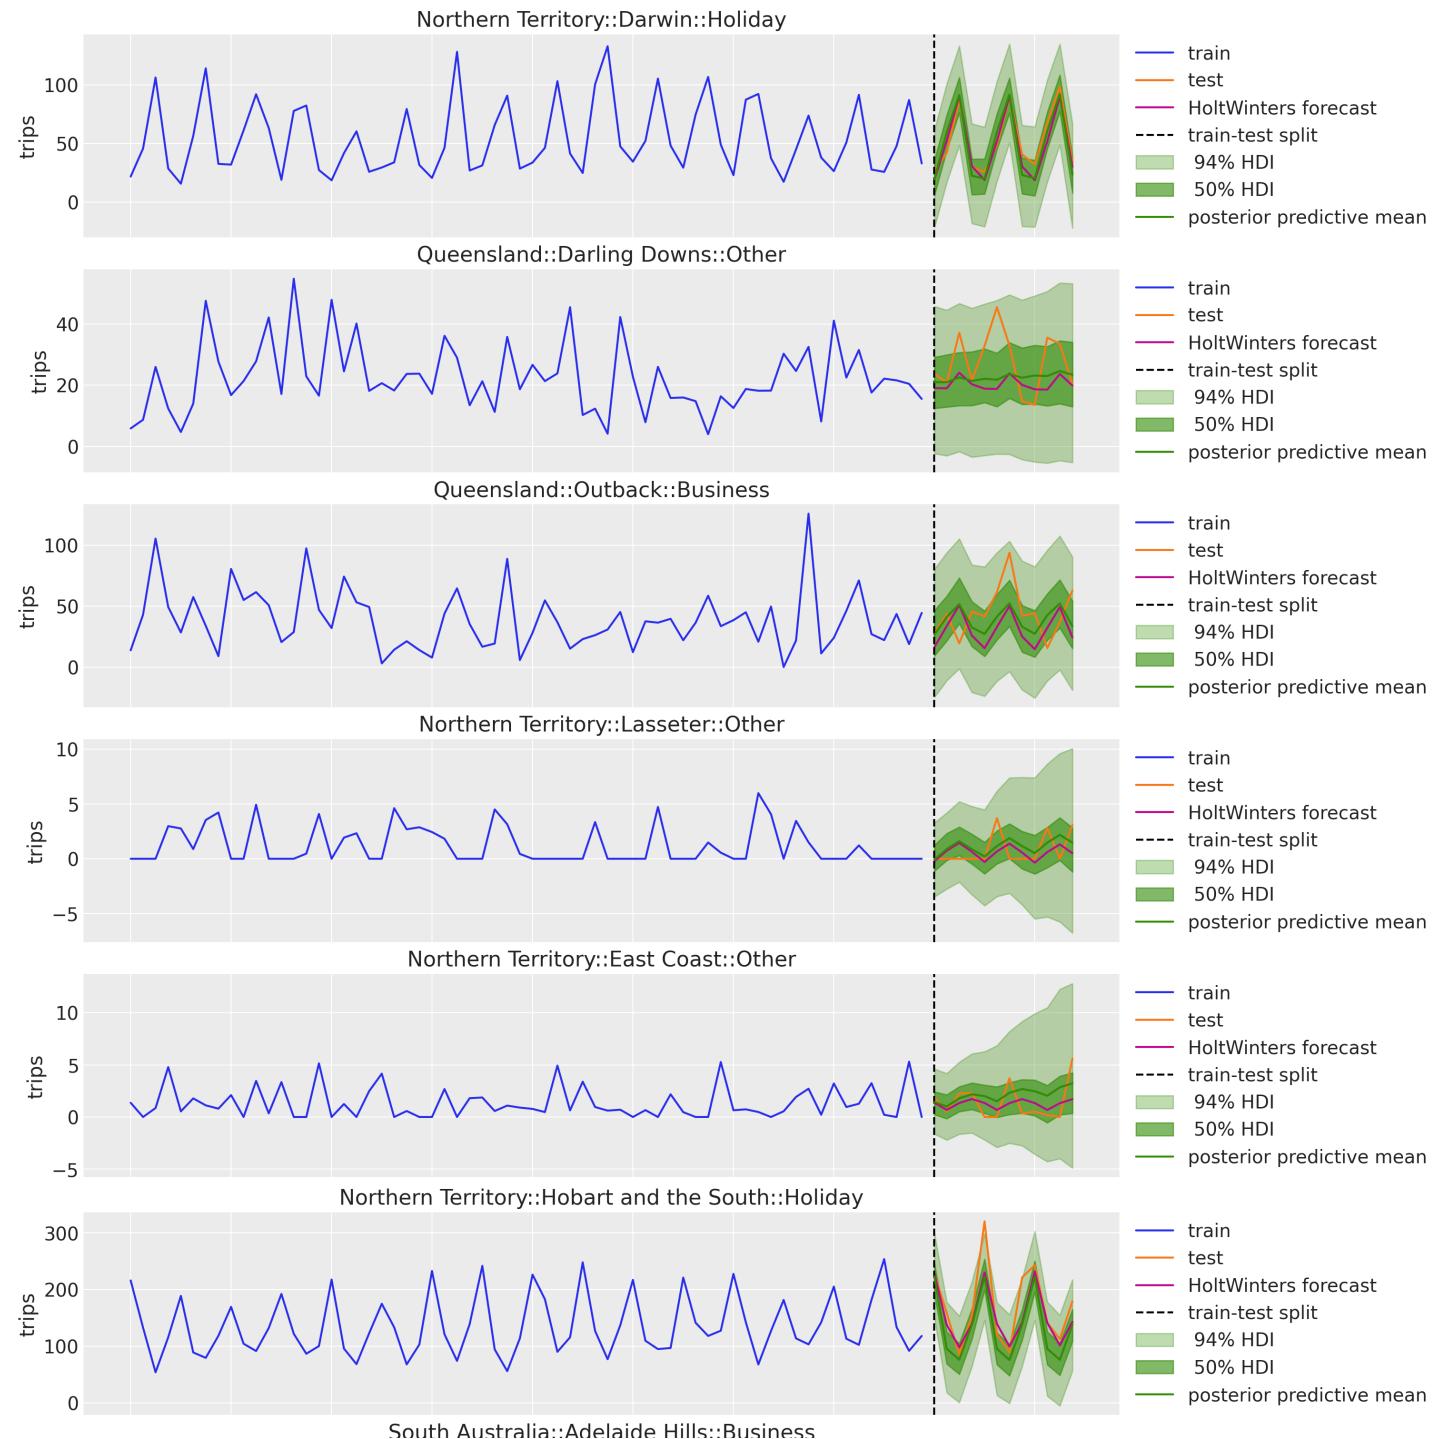
```

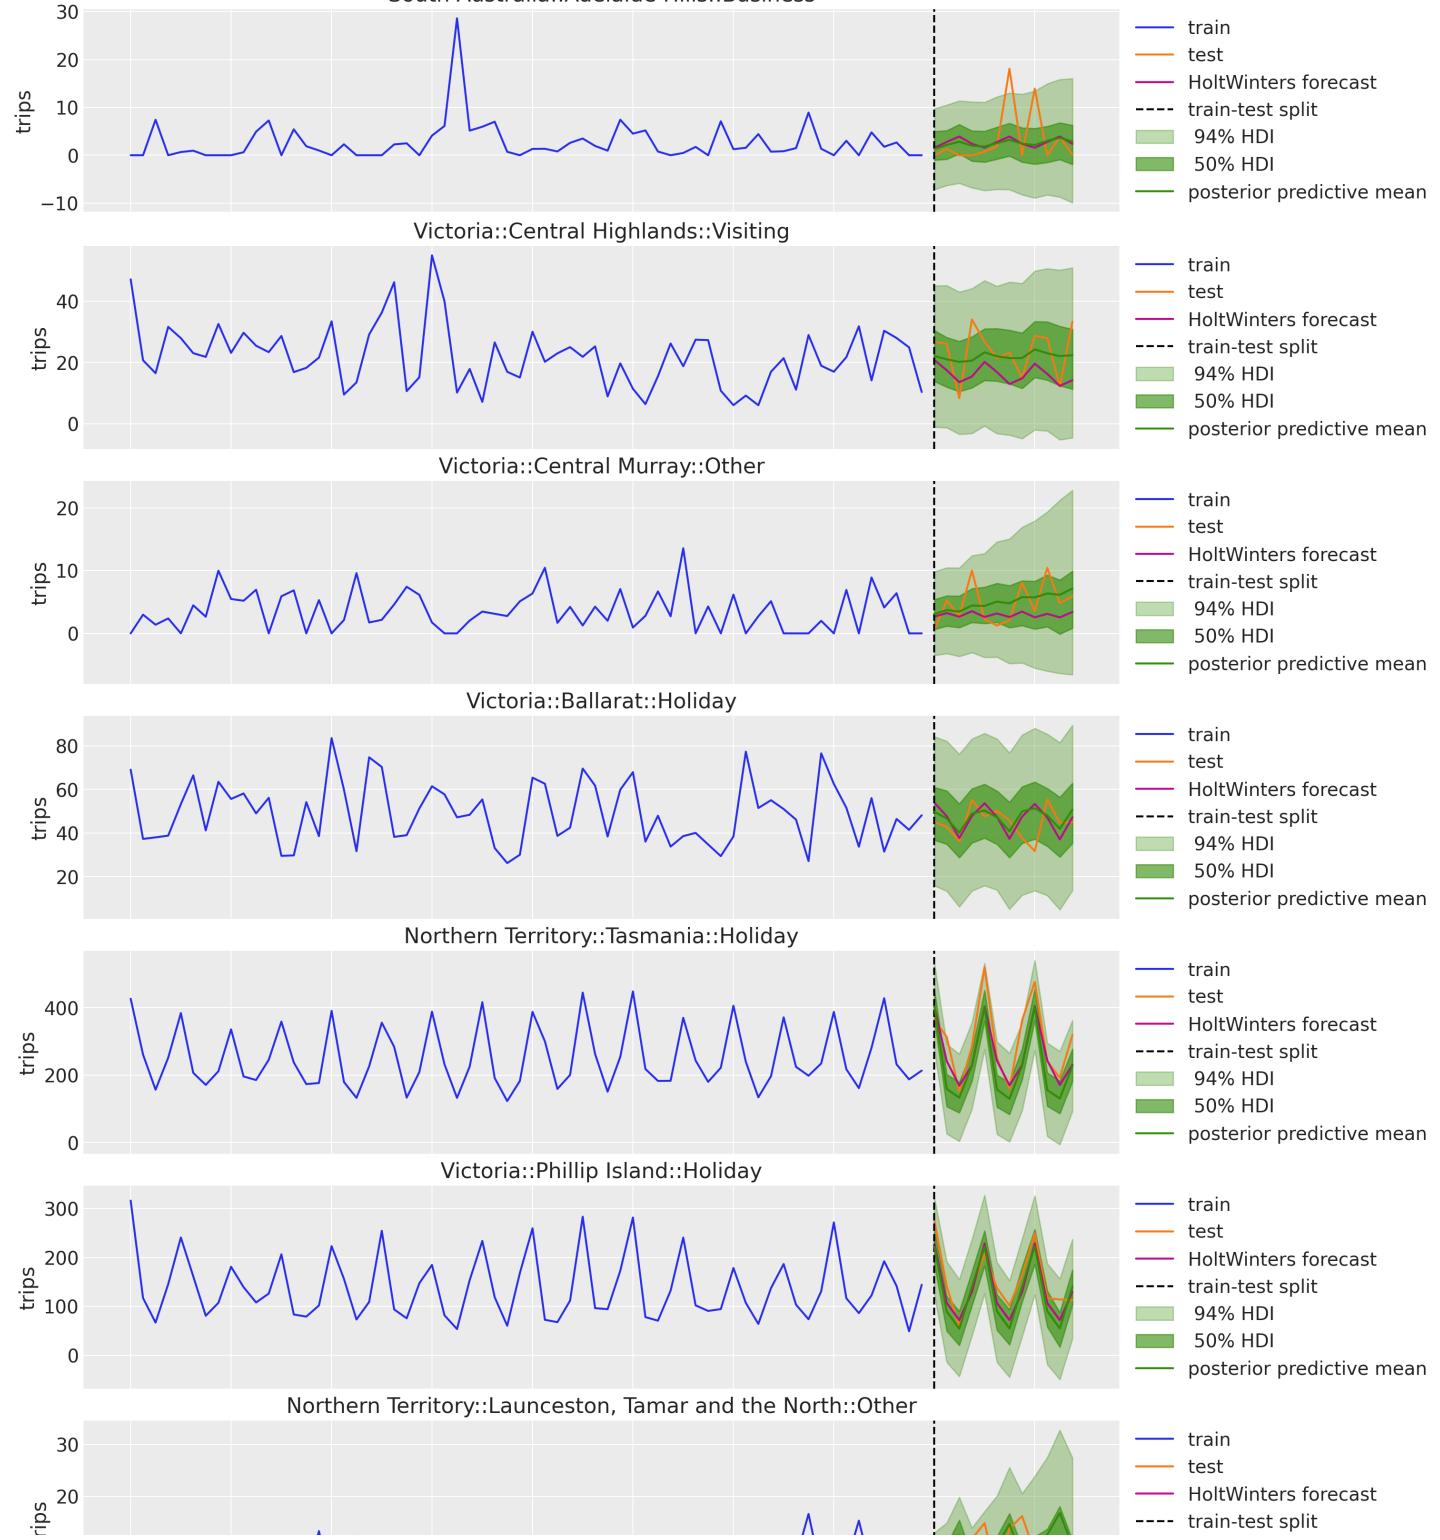


Hierarchical Exponential Smoothing

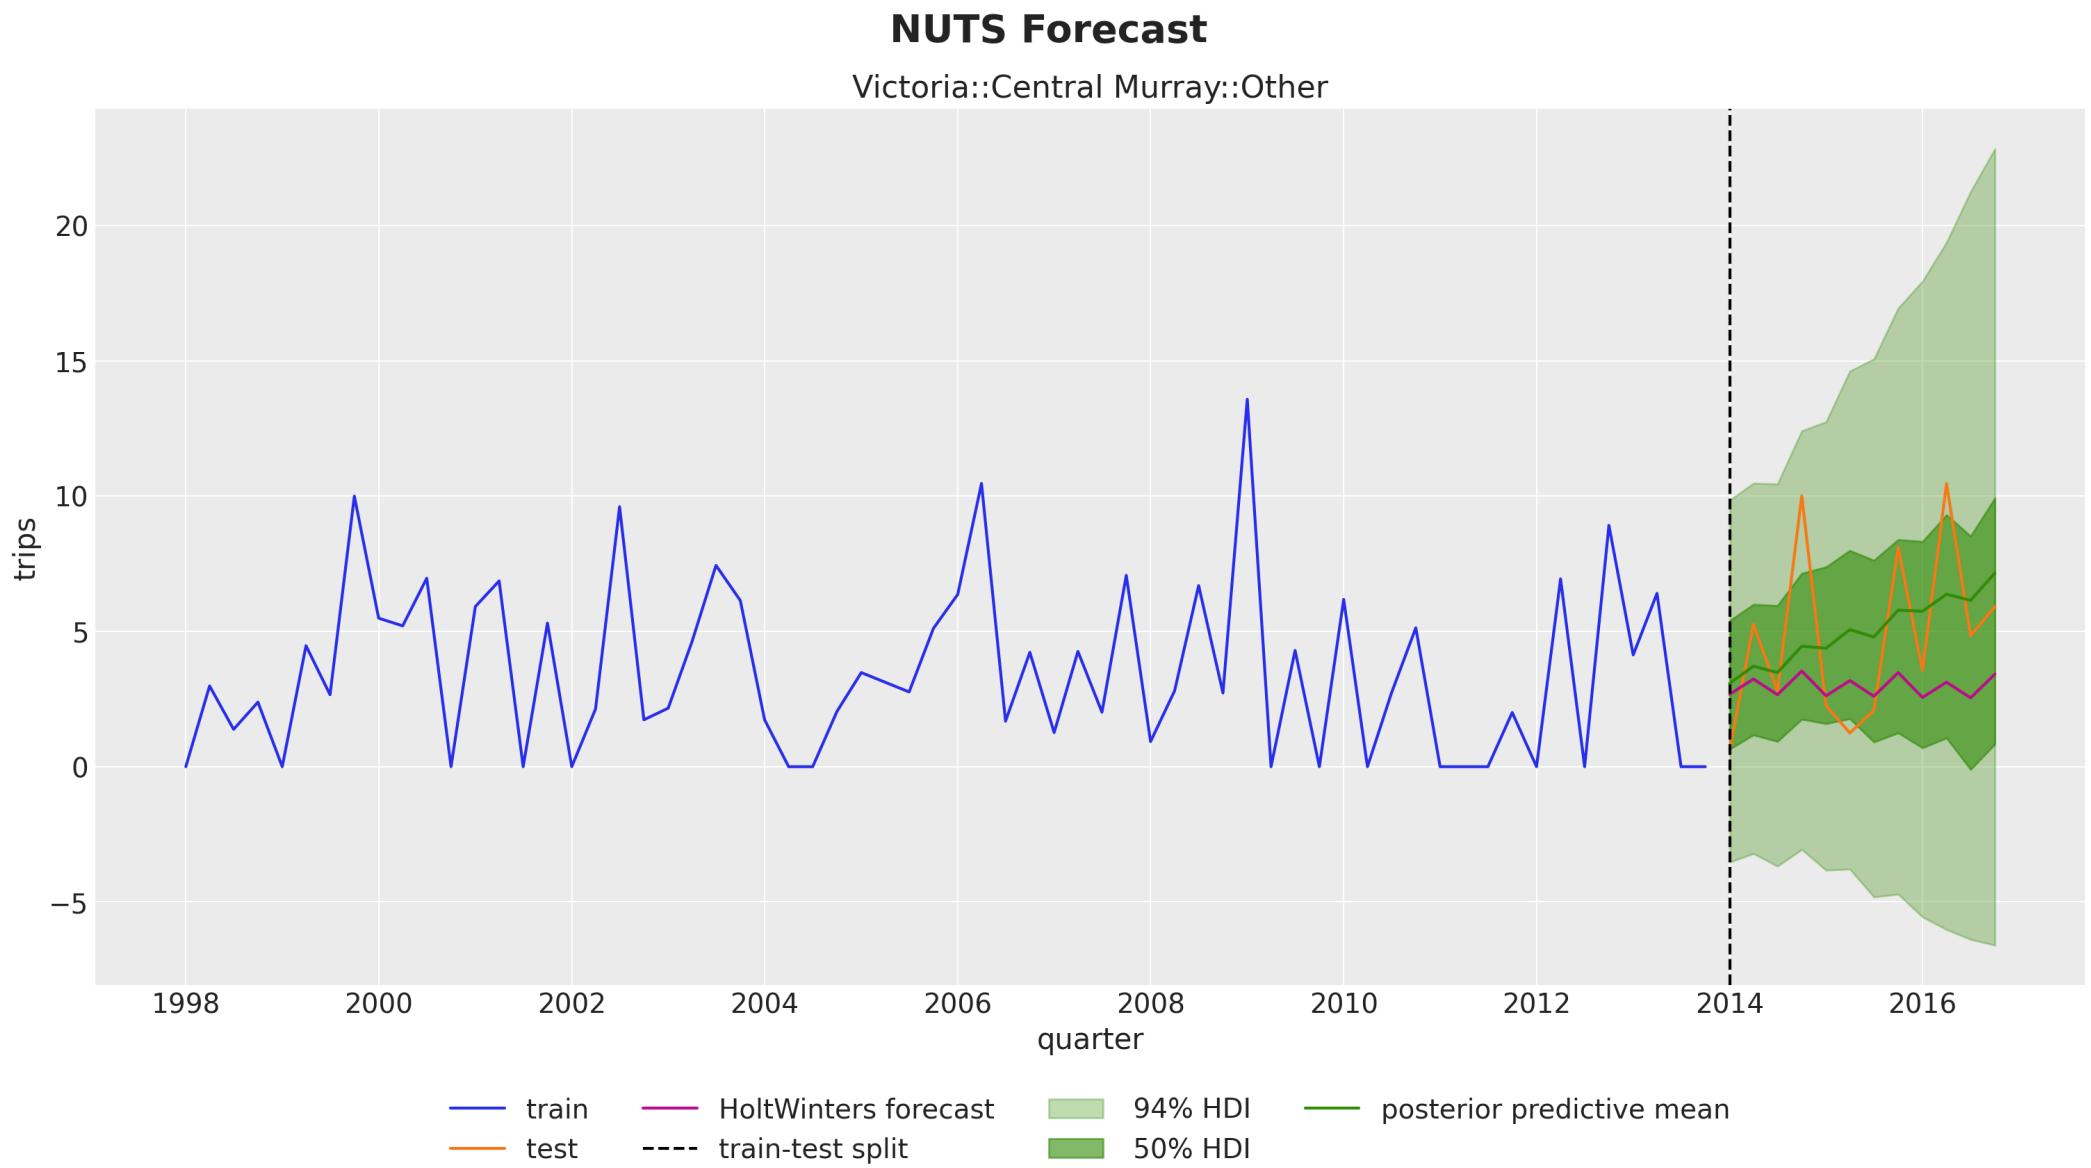


NUTS Forecast

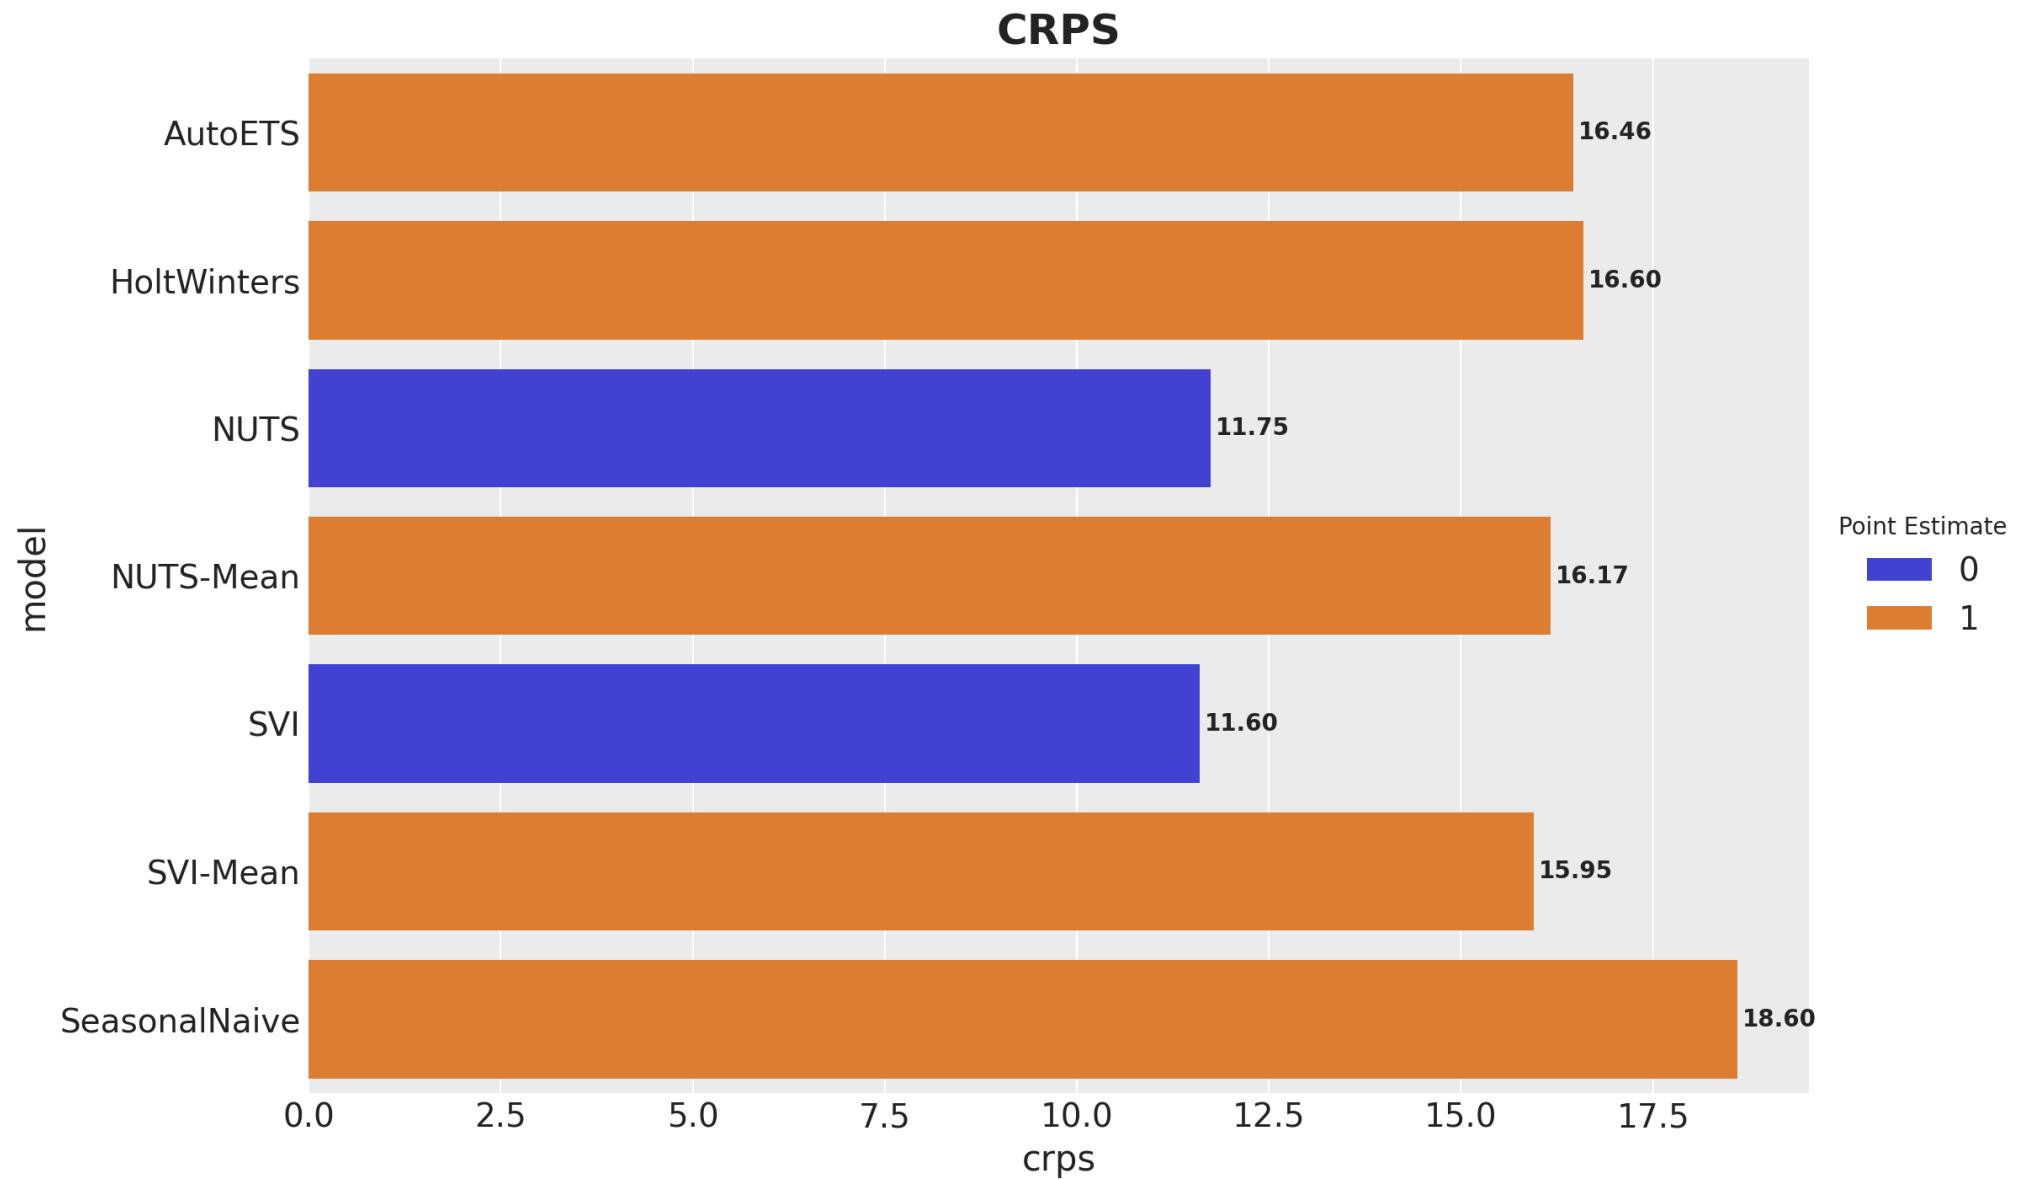




Hierarchical Exponential Smoothing

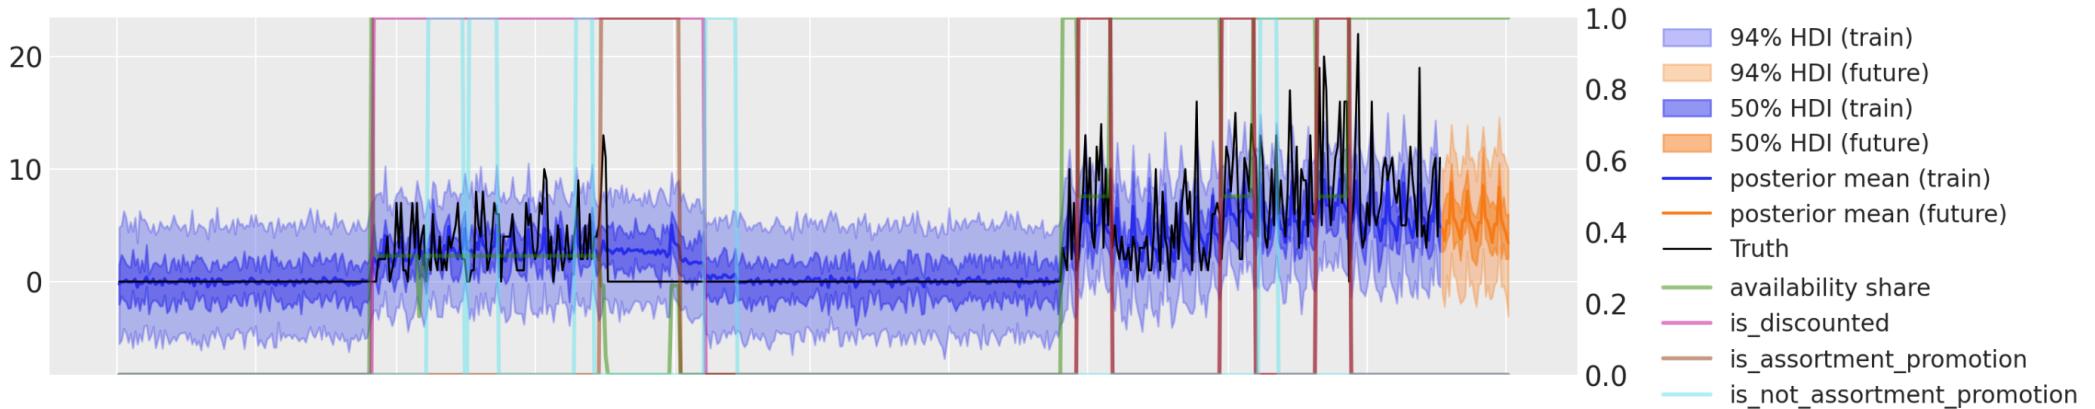


Hierarchical Exponential Smoothing



Baseline Model

Local Level Model + Seasonality + Covariates

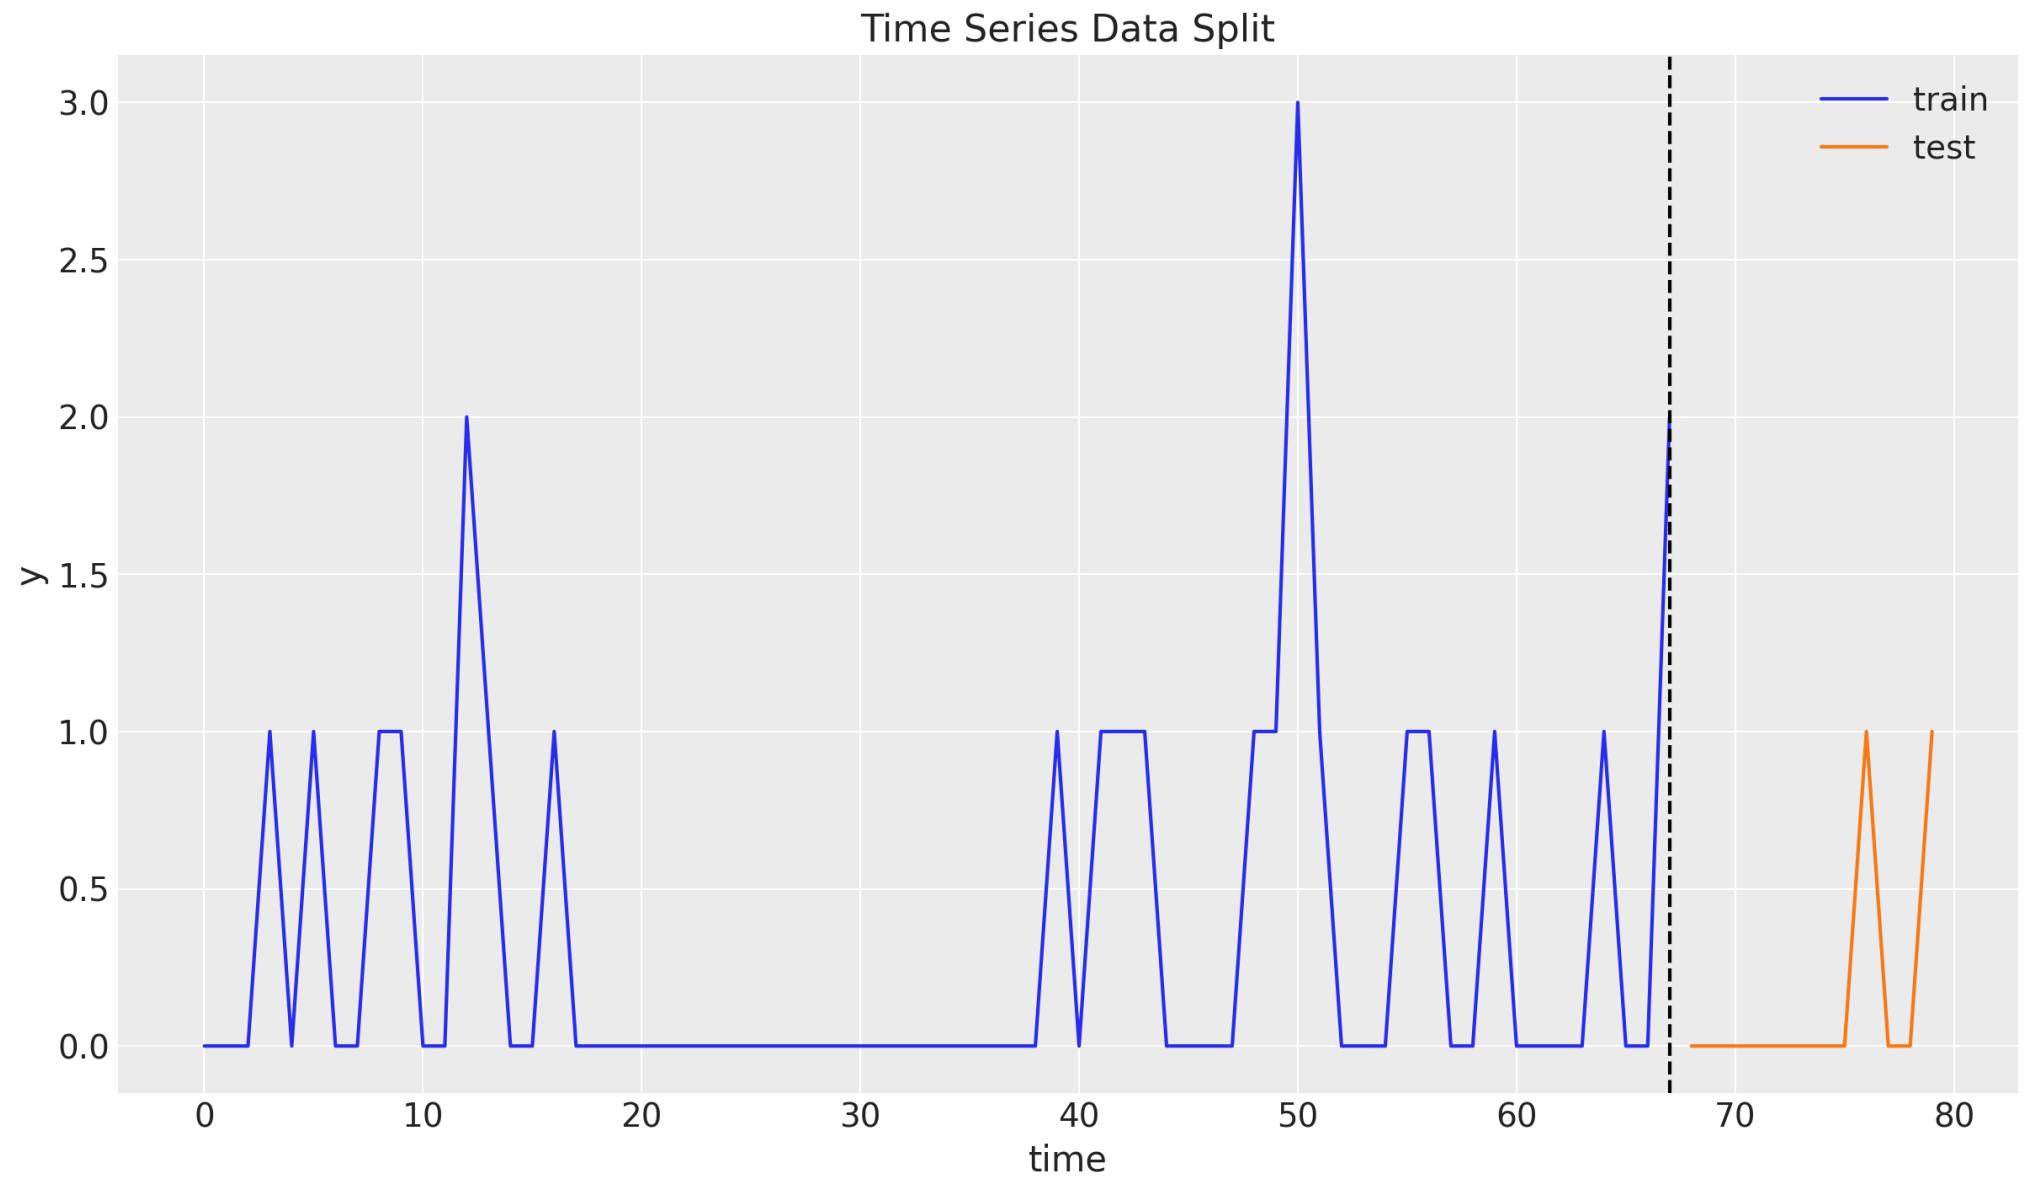


- Local level model to capture the trend component.
- Seasonality using a Fourier modes.
- Add covariates to account for promotions and discounts.
- Global factor to account the availability of the product.

Scalability: 40K time-series can be fitted in less than 10 minutes in a GPU.



Intermittent Time Series



Croston's Method

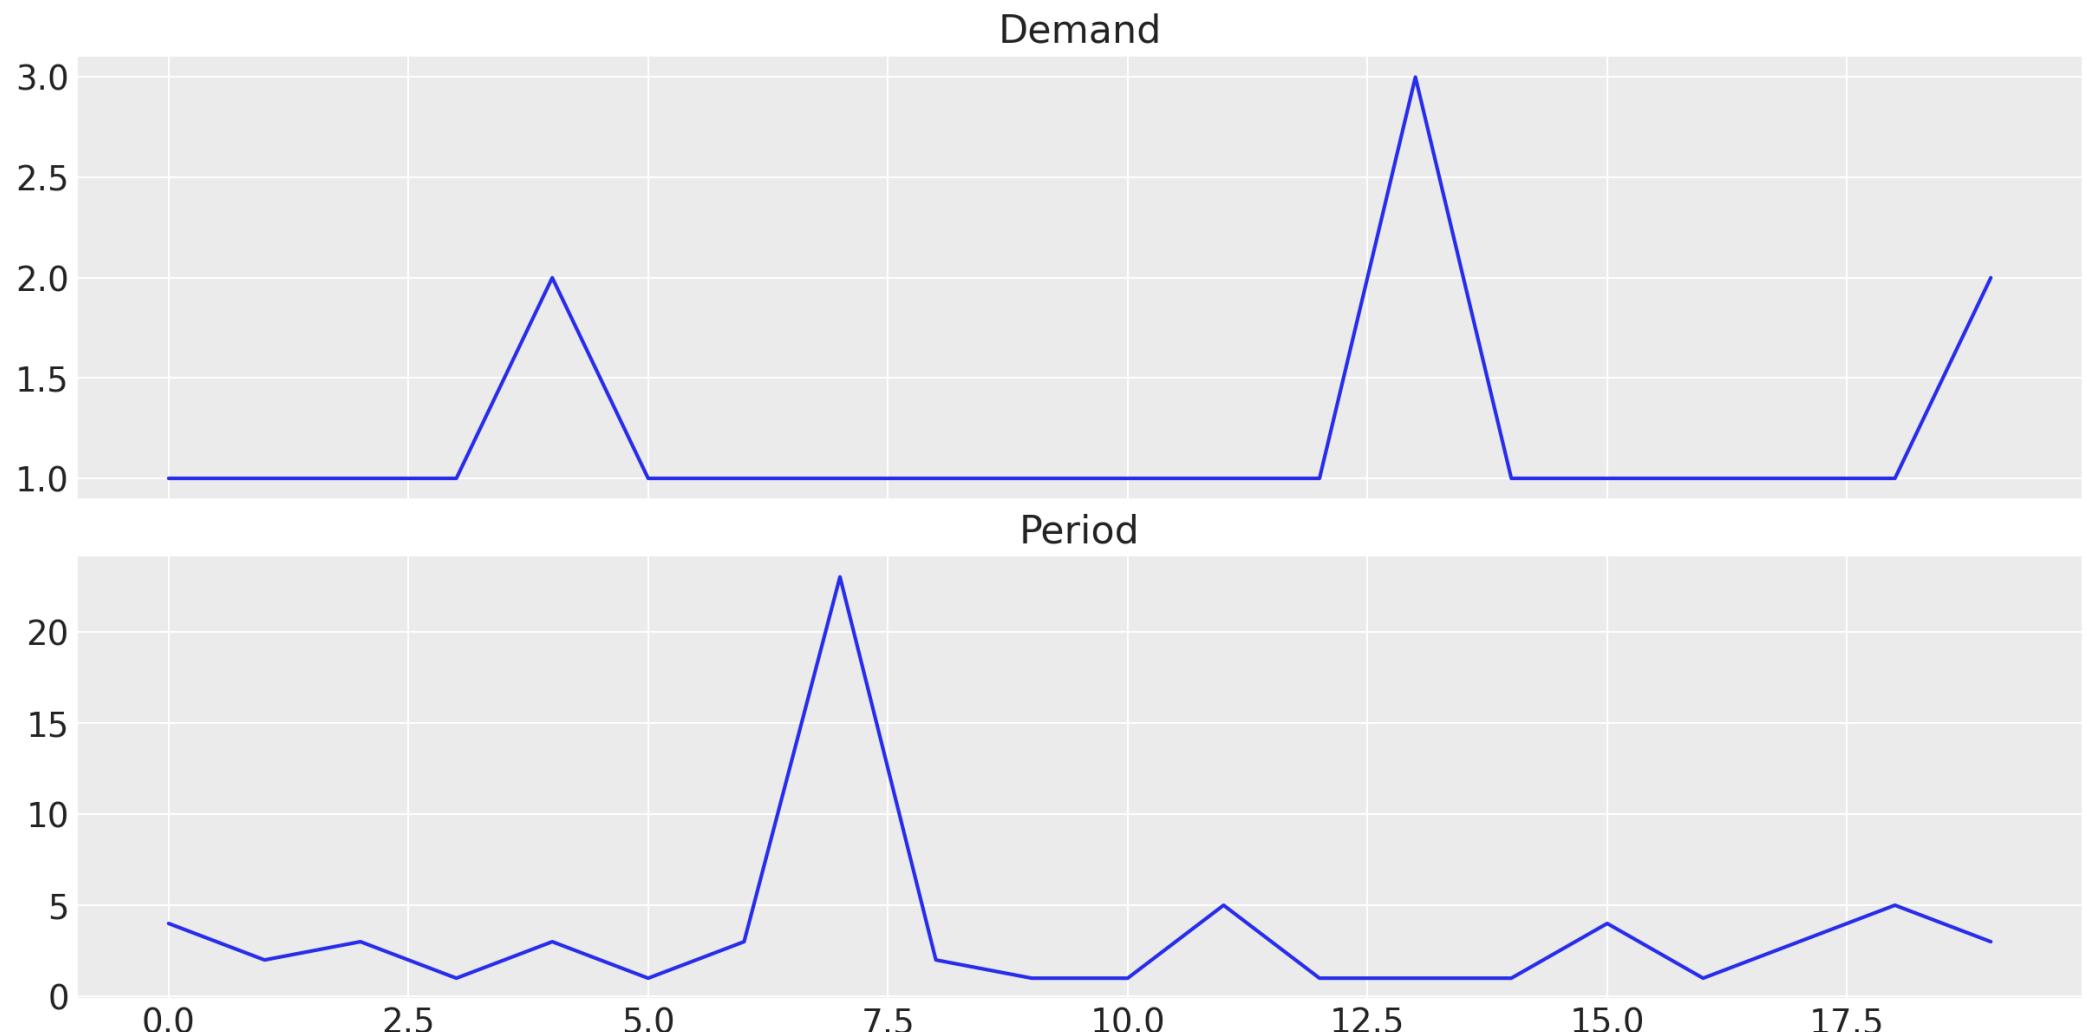
 The method is based on the idea of **separating the demand size z_t and the demand interval p_t** , and then **forecasting them separately** using simple exponential smoothing.

- z_t : keep the non-zero values of y_t .
- p_t : keep the time between non-zero values of y_t .

$$\hat{y}_{t+h} = \frac{\hat{z}_{t+h}}{\hat{p}_{t+h}}$$



Croston's Method

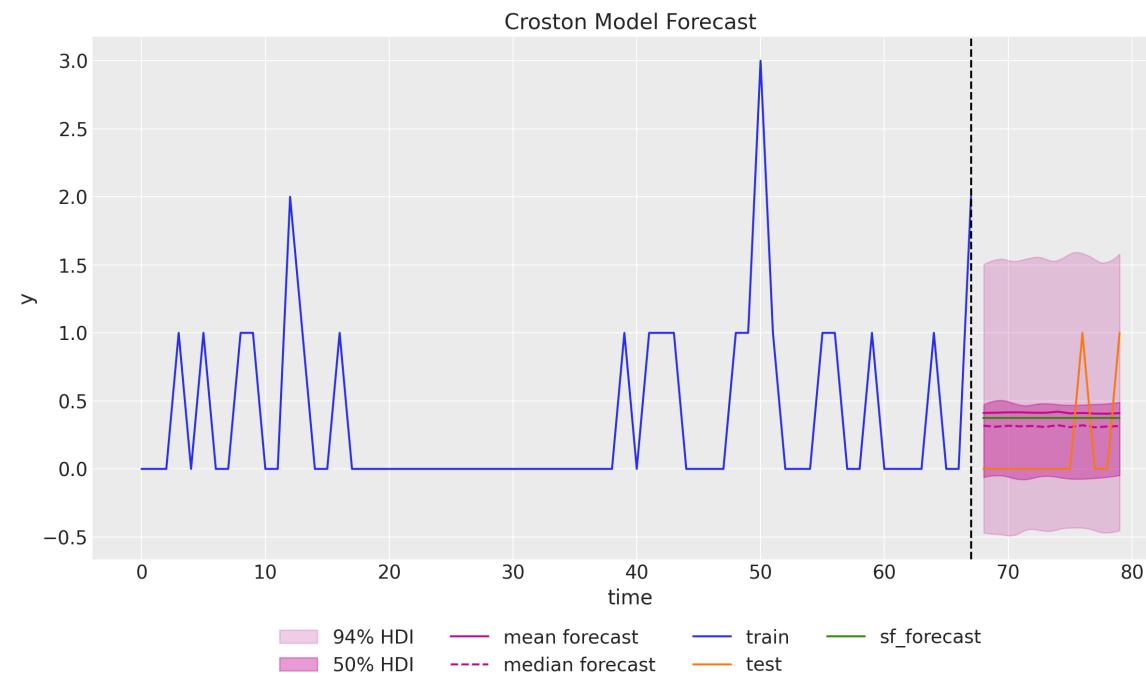


Croston's Method

```

1 def croston_model(z: ArrayLike, p_inv: ArrayLike, future: int = 0) -> None:
2     z_forecast = scope(level_model, "demand")(z, future)
3     p_inv_forecast = scope(level_model, "period_inv")(p_inv, future)
4
5     if future > 0:
6         numpyro.deterministic("z_forecast", z_forecast)
7         numpyro.deterministic("p_inv_forecast", p_inv_forecast)
8         numpyro.deterministic("forecast", z_forecast * p_inv_forecast)

```



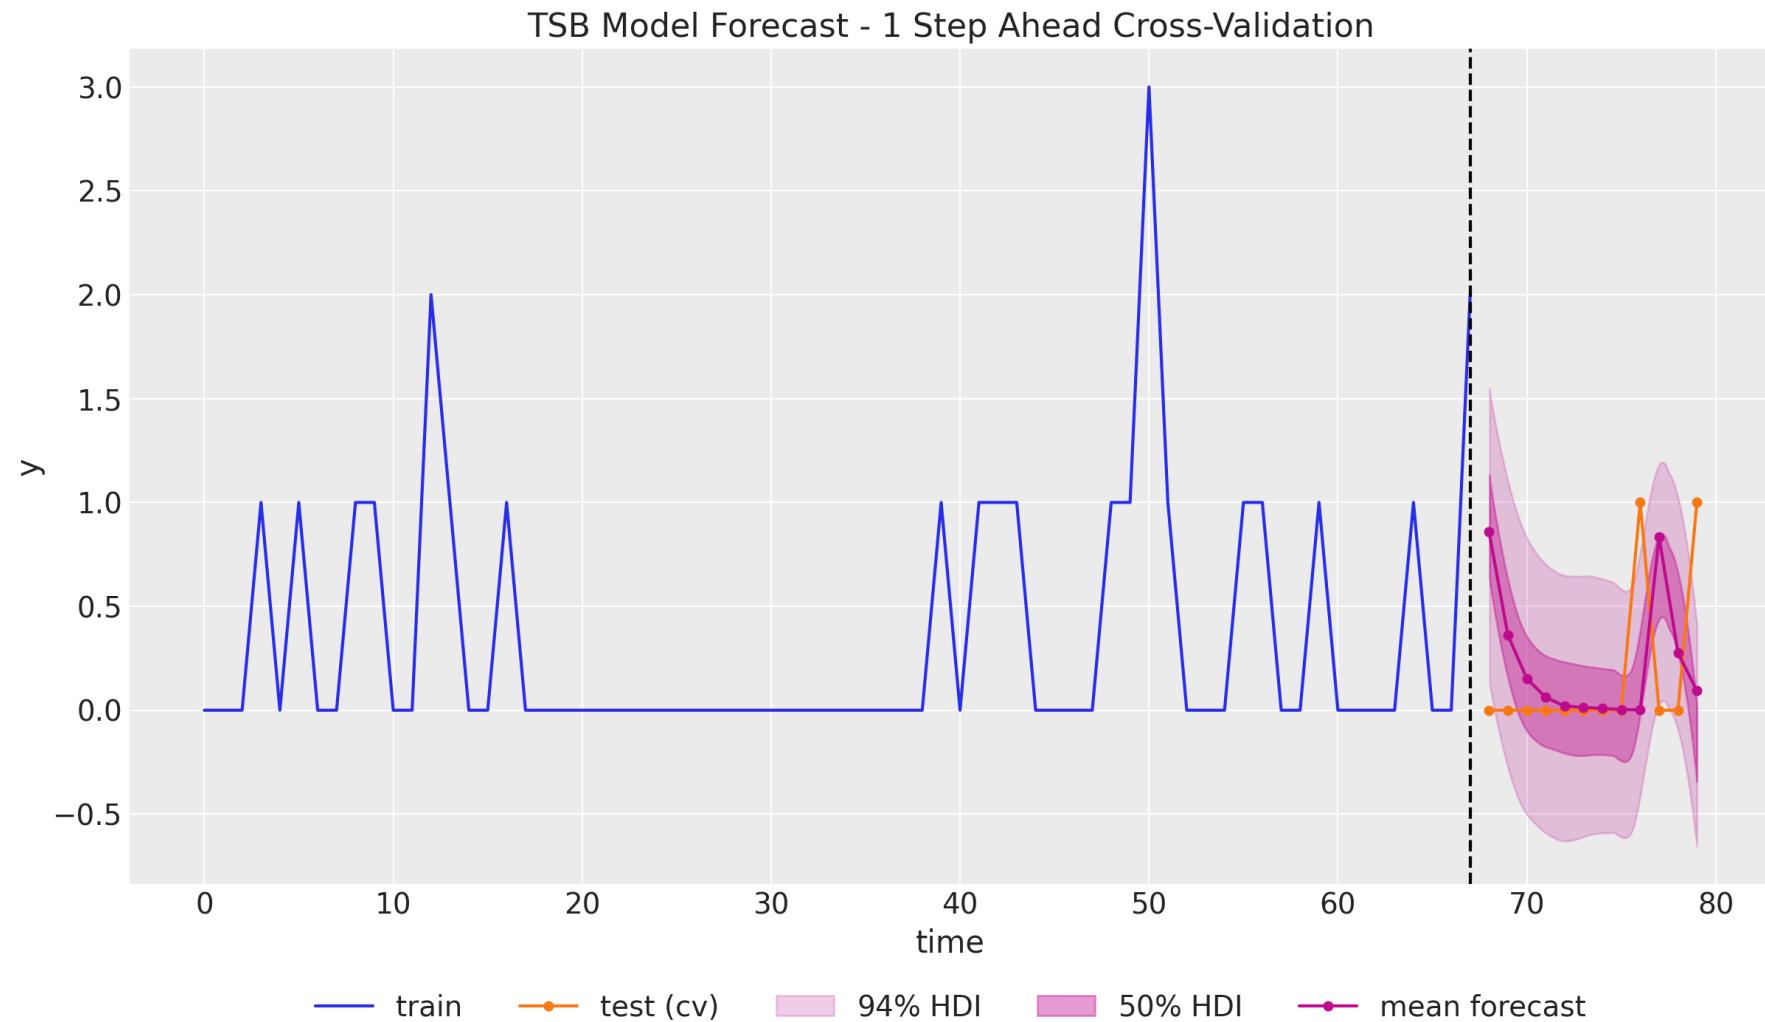
TSB Method

- The TSB method is similar to the Croston's method: constructs two different time series out of the original one and then forecast each of them separately, so that the final forecast is generated by combining the forecasts of the two time series.
- **The main difference between the two methods is that the TSB method uses the demand probability instead of the demand periods.**



TSB Method

1 - Step Ahead Time Slice Cross Validation



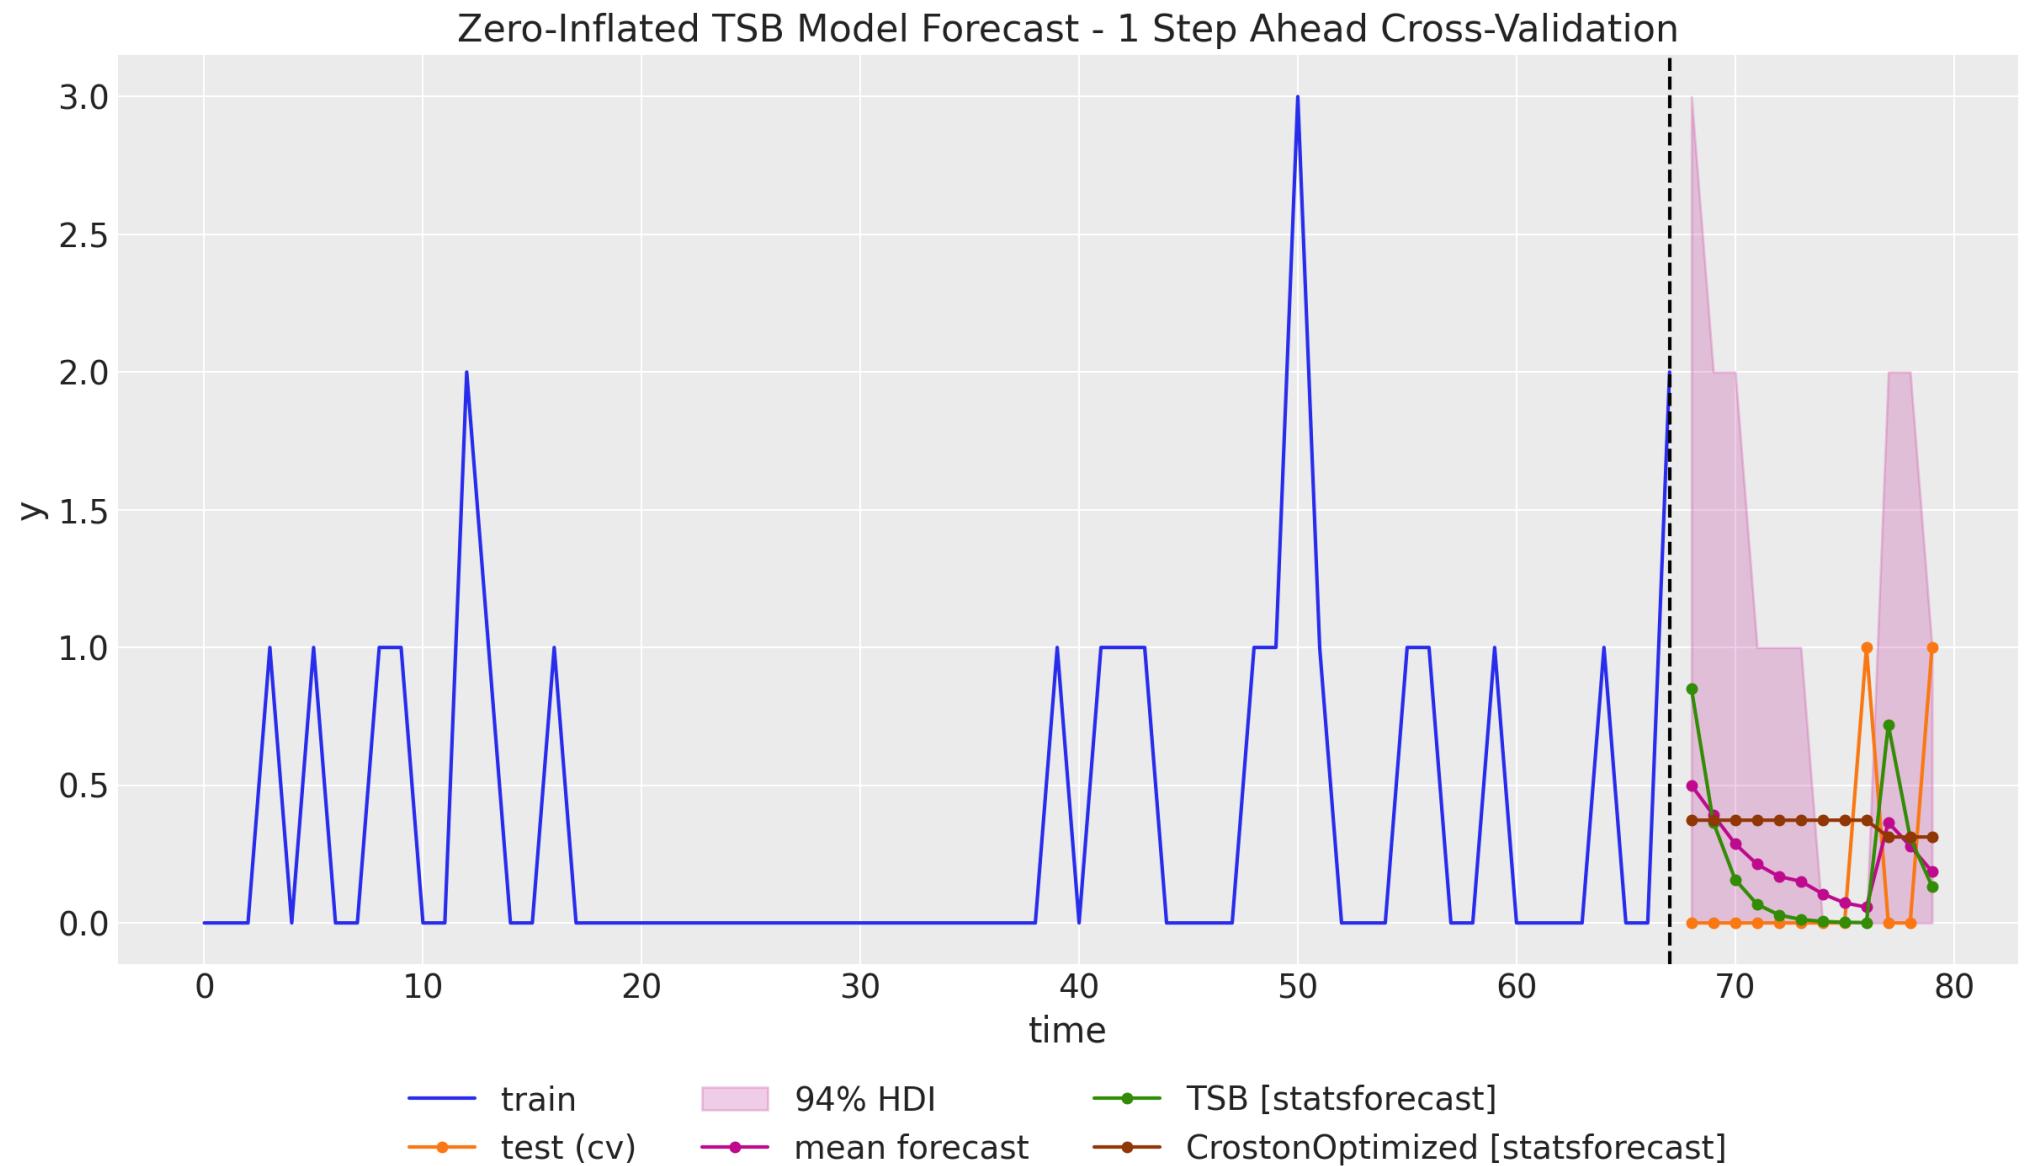
Zero-Inflated TSB Model

🧪 We can modify the TSB model to include zero-inflation by using a Zero-Inflated Negative Binomial Distribution.

```
1  def transition_fn(carry, t):
2      z_prev, p_prev = carry
3
4      z_next . . .
5      p_next . . .
6
7      mu = z_next
8      gate = 1 - p_next
9      pred = numpyro.sample(
10          "pred",
11          dist.ZeroInflatedNegativeBinomial2(
12              mean=mu, concentration=concentration, gate=gate
13          ),
14      )
```

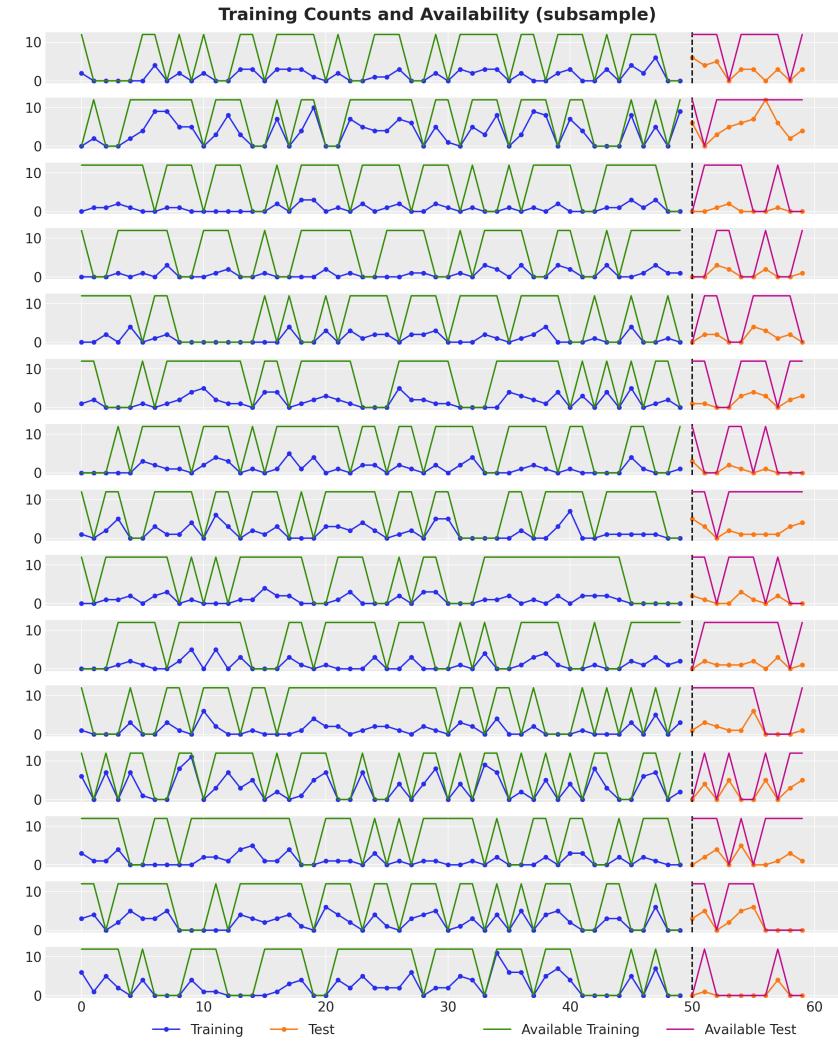


Time-Slice Cross Validation



Why Zeros Happen?

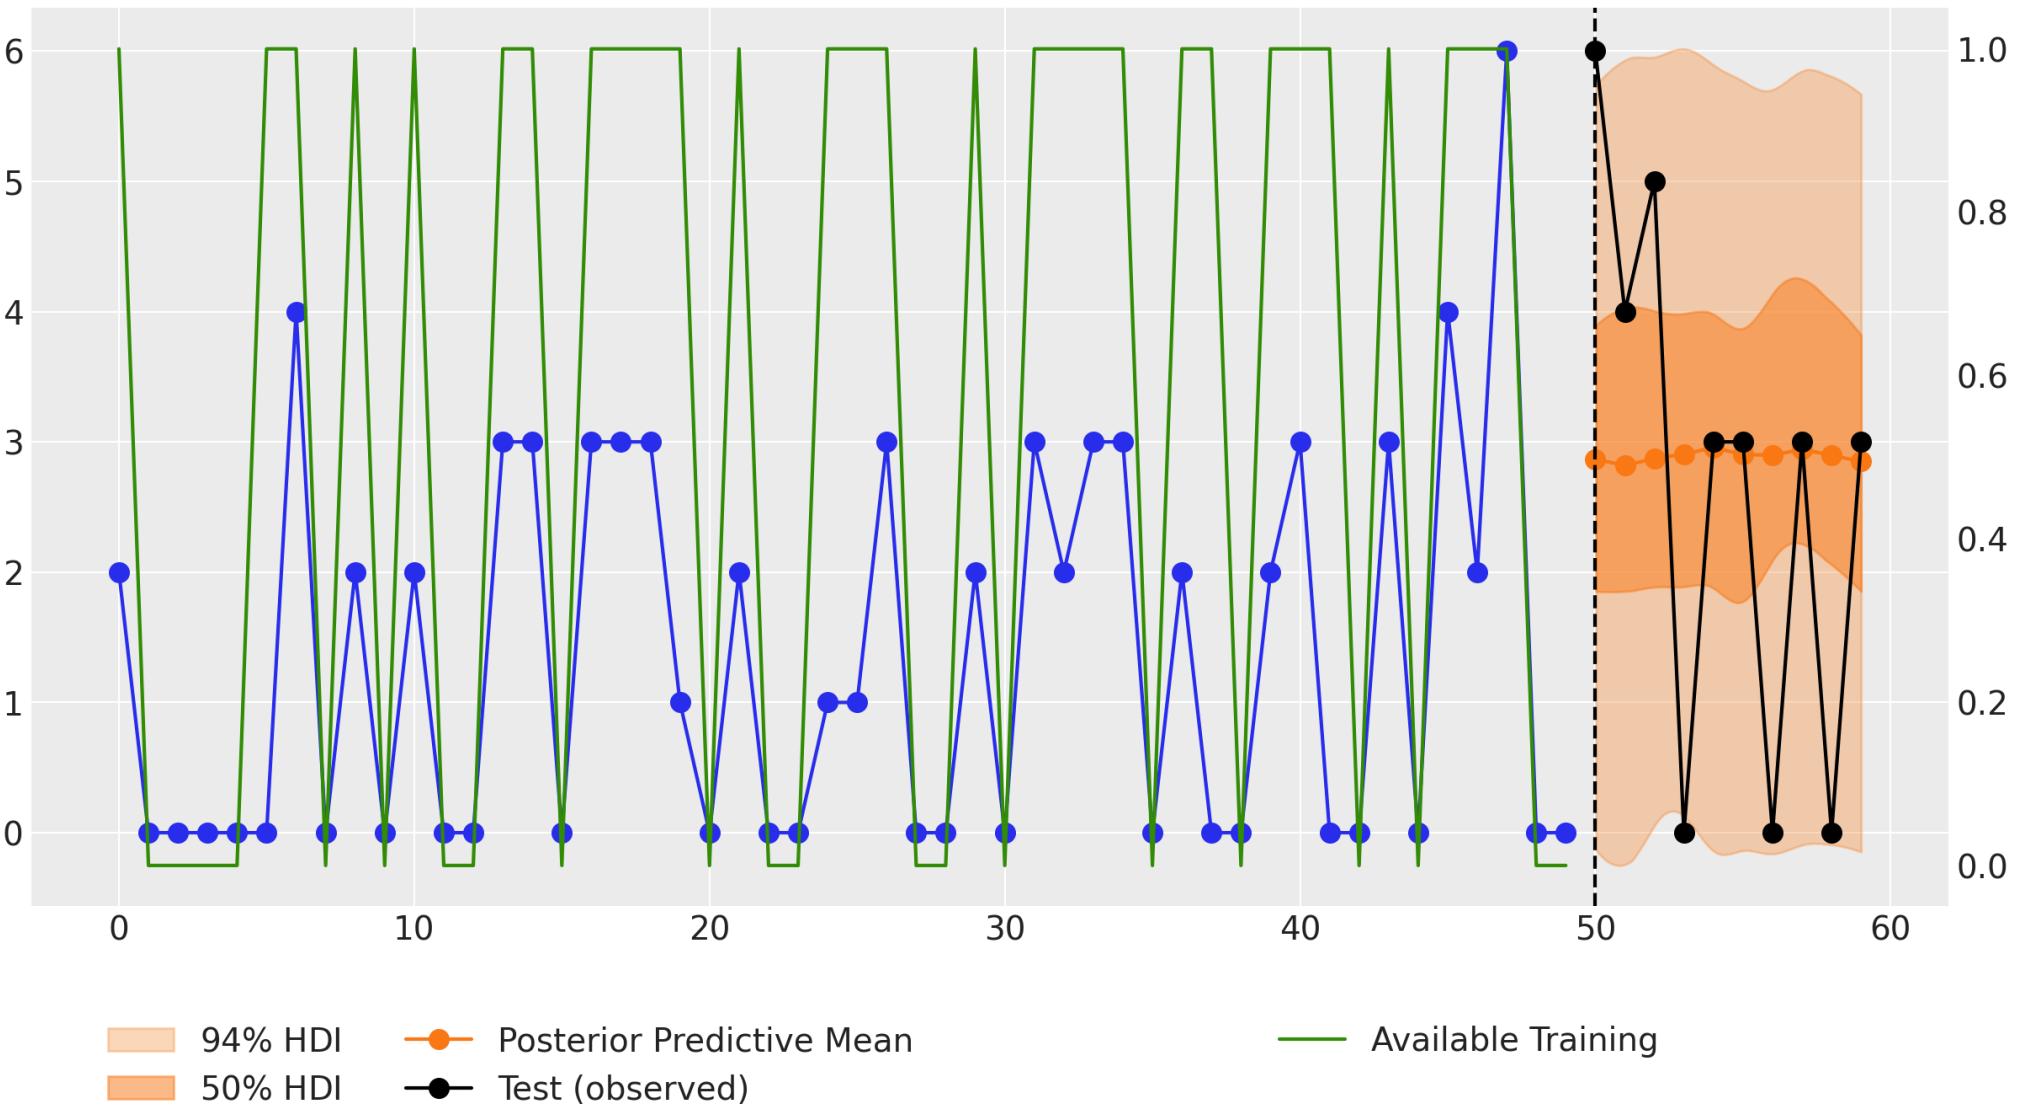
Simulation Study: Availability Constraints



Hacking the TSB Model

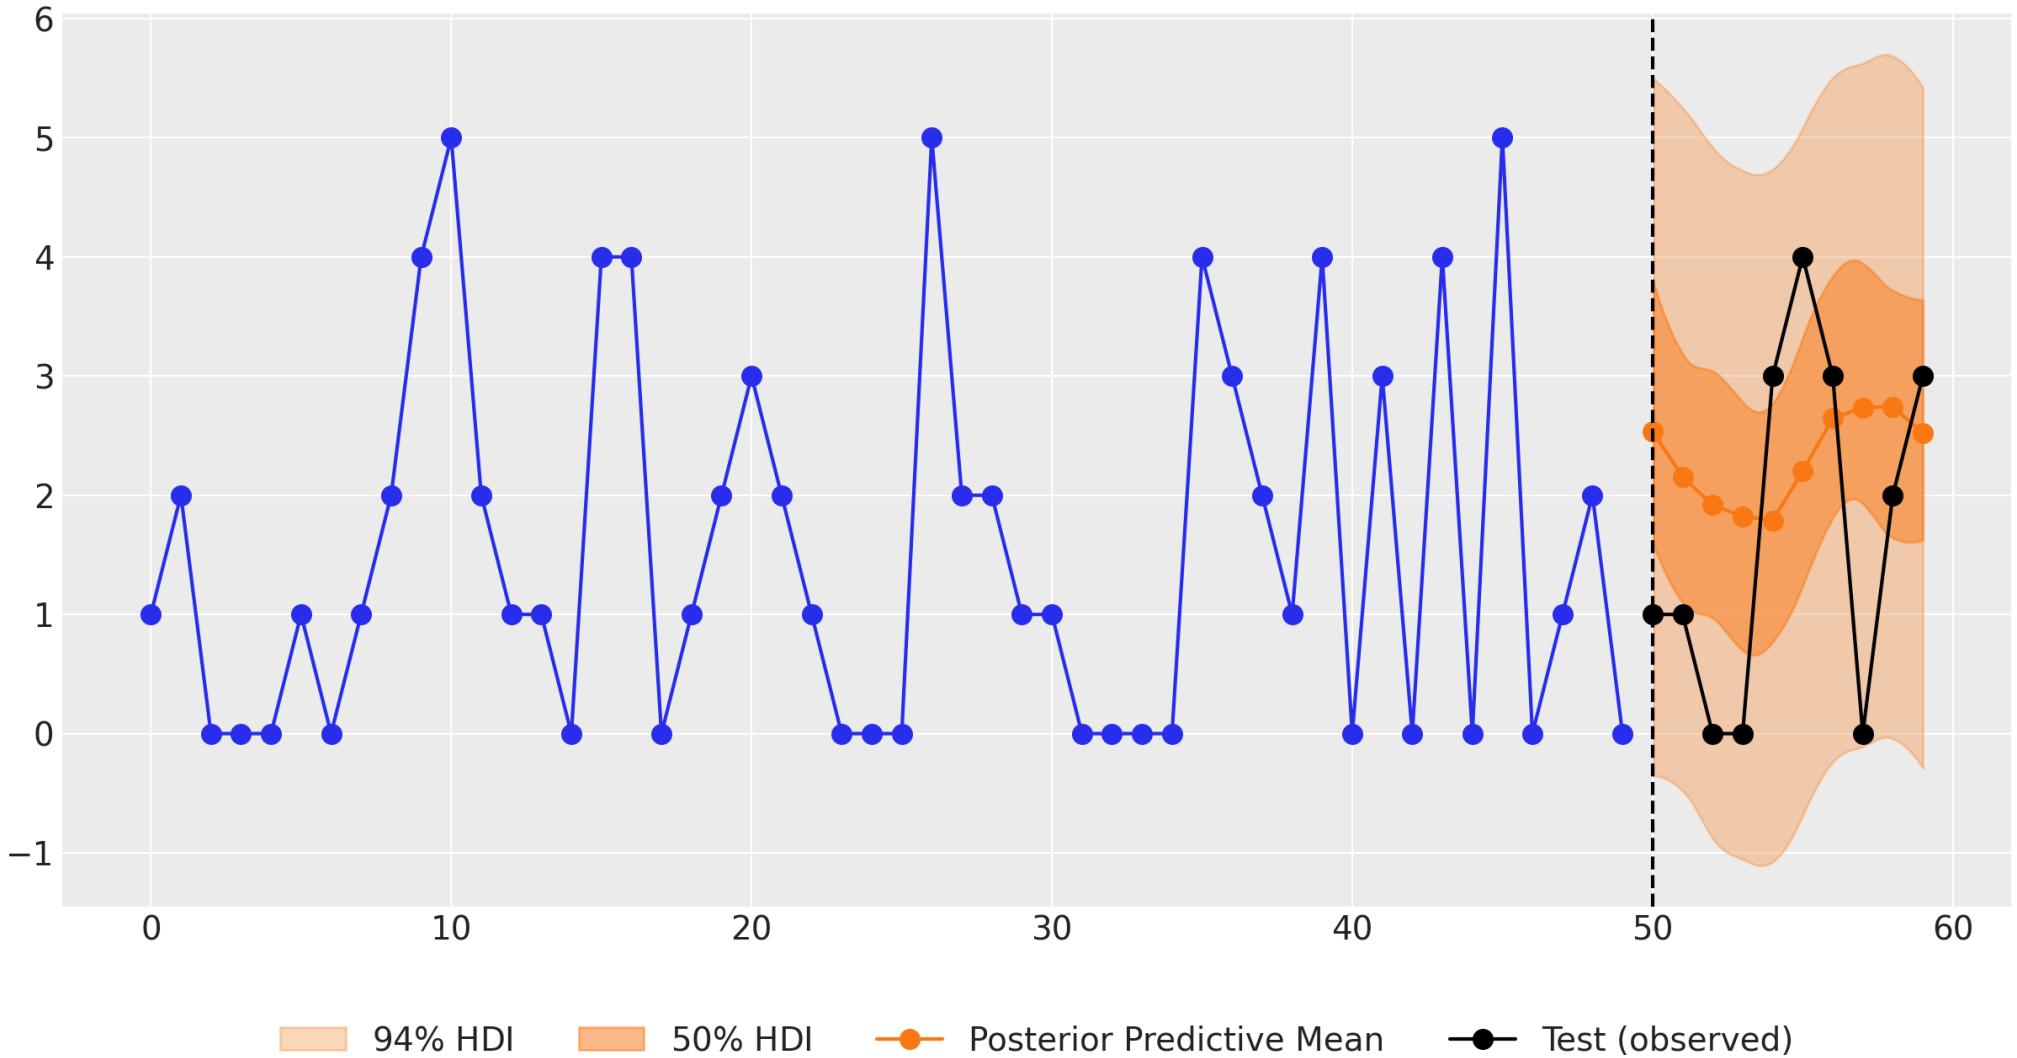


Posterior Predictive Checks



Hacking the TSB Model

**Posterior Predictive
1 Step Ahead Cross-Validation**

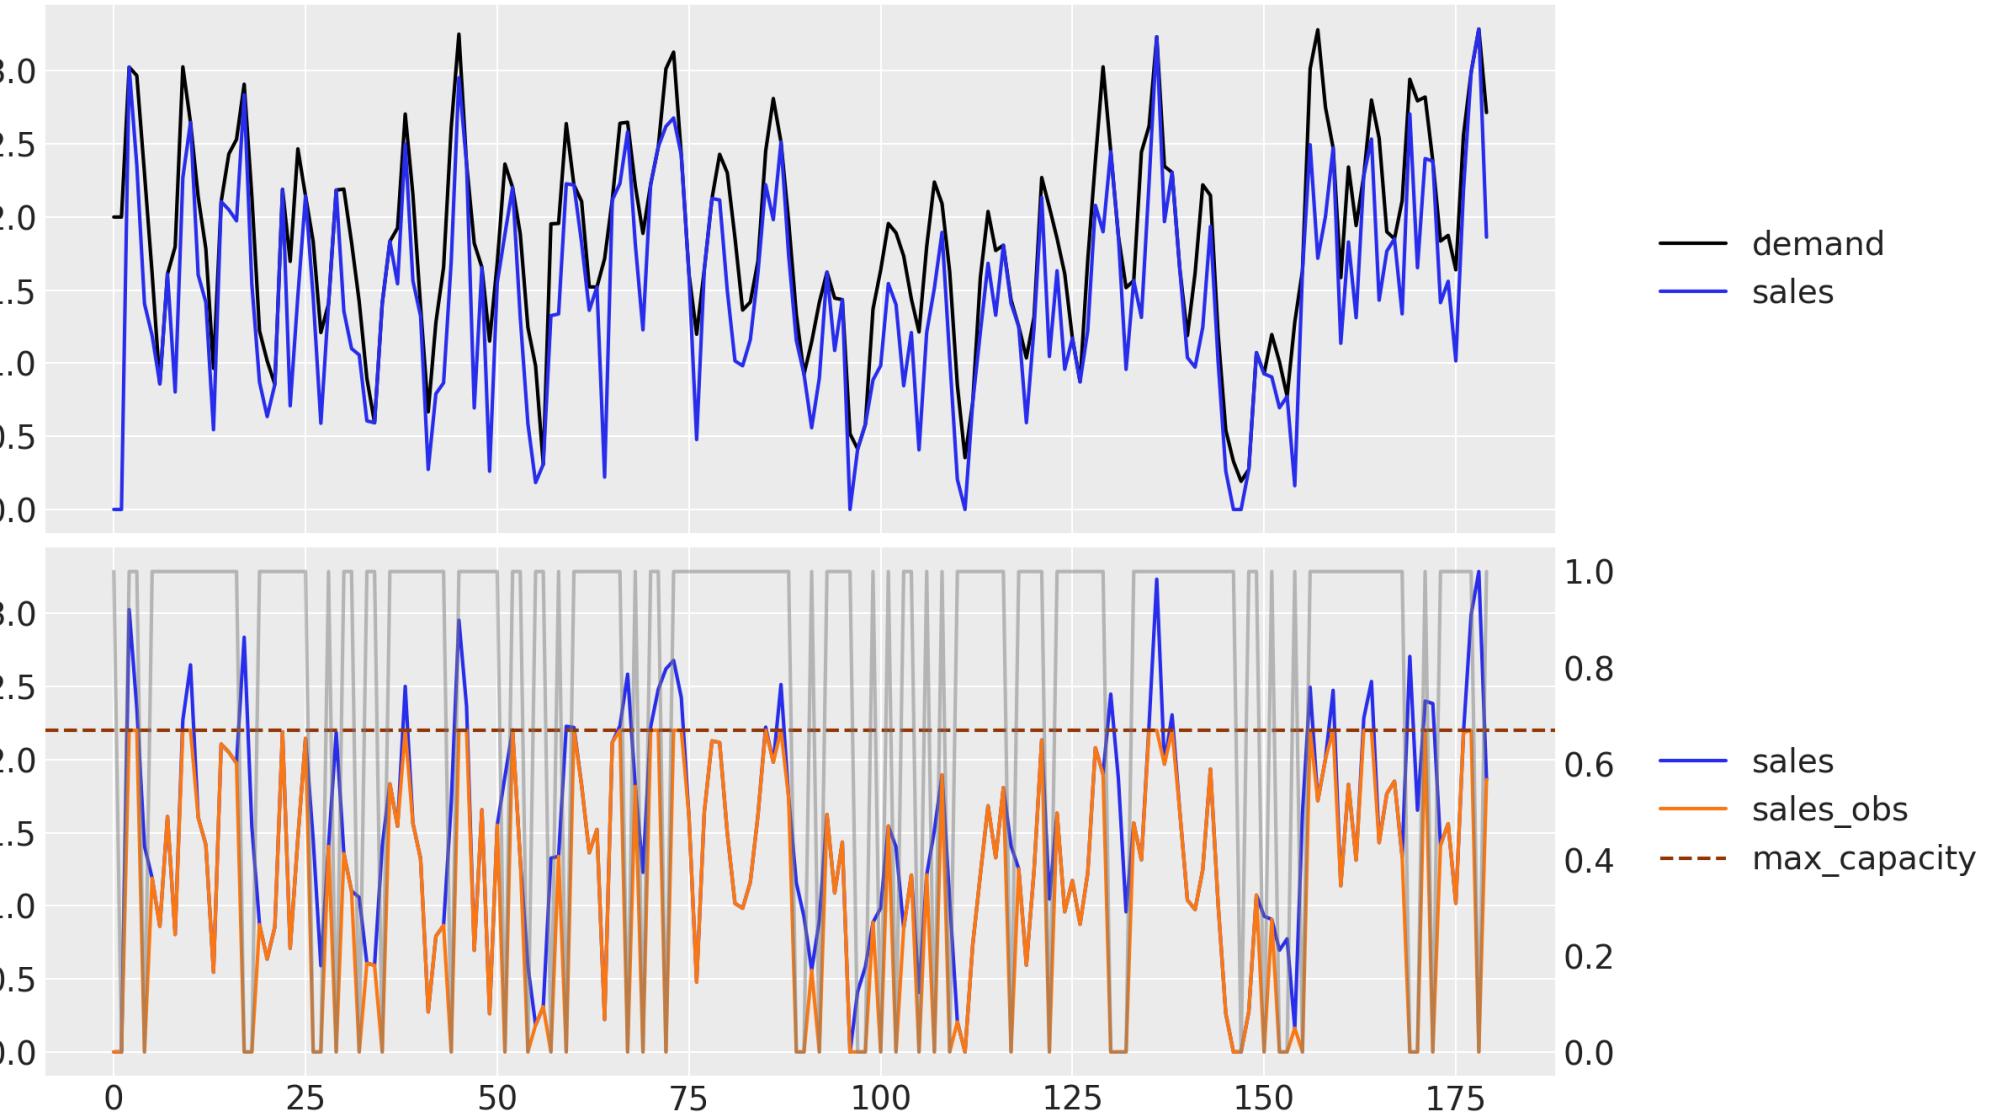


■ 94% HDI ■ 50% HDI ■● Posterior Predictive Mean ■● Test (observed)



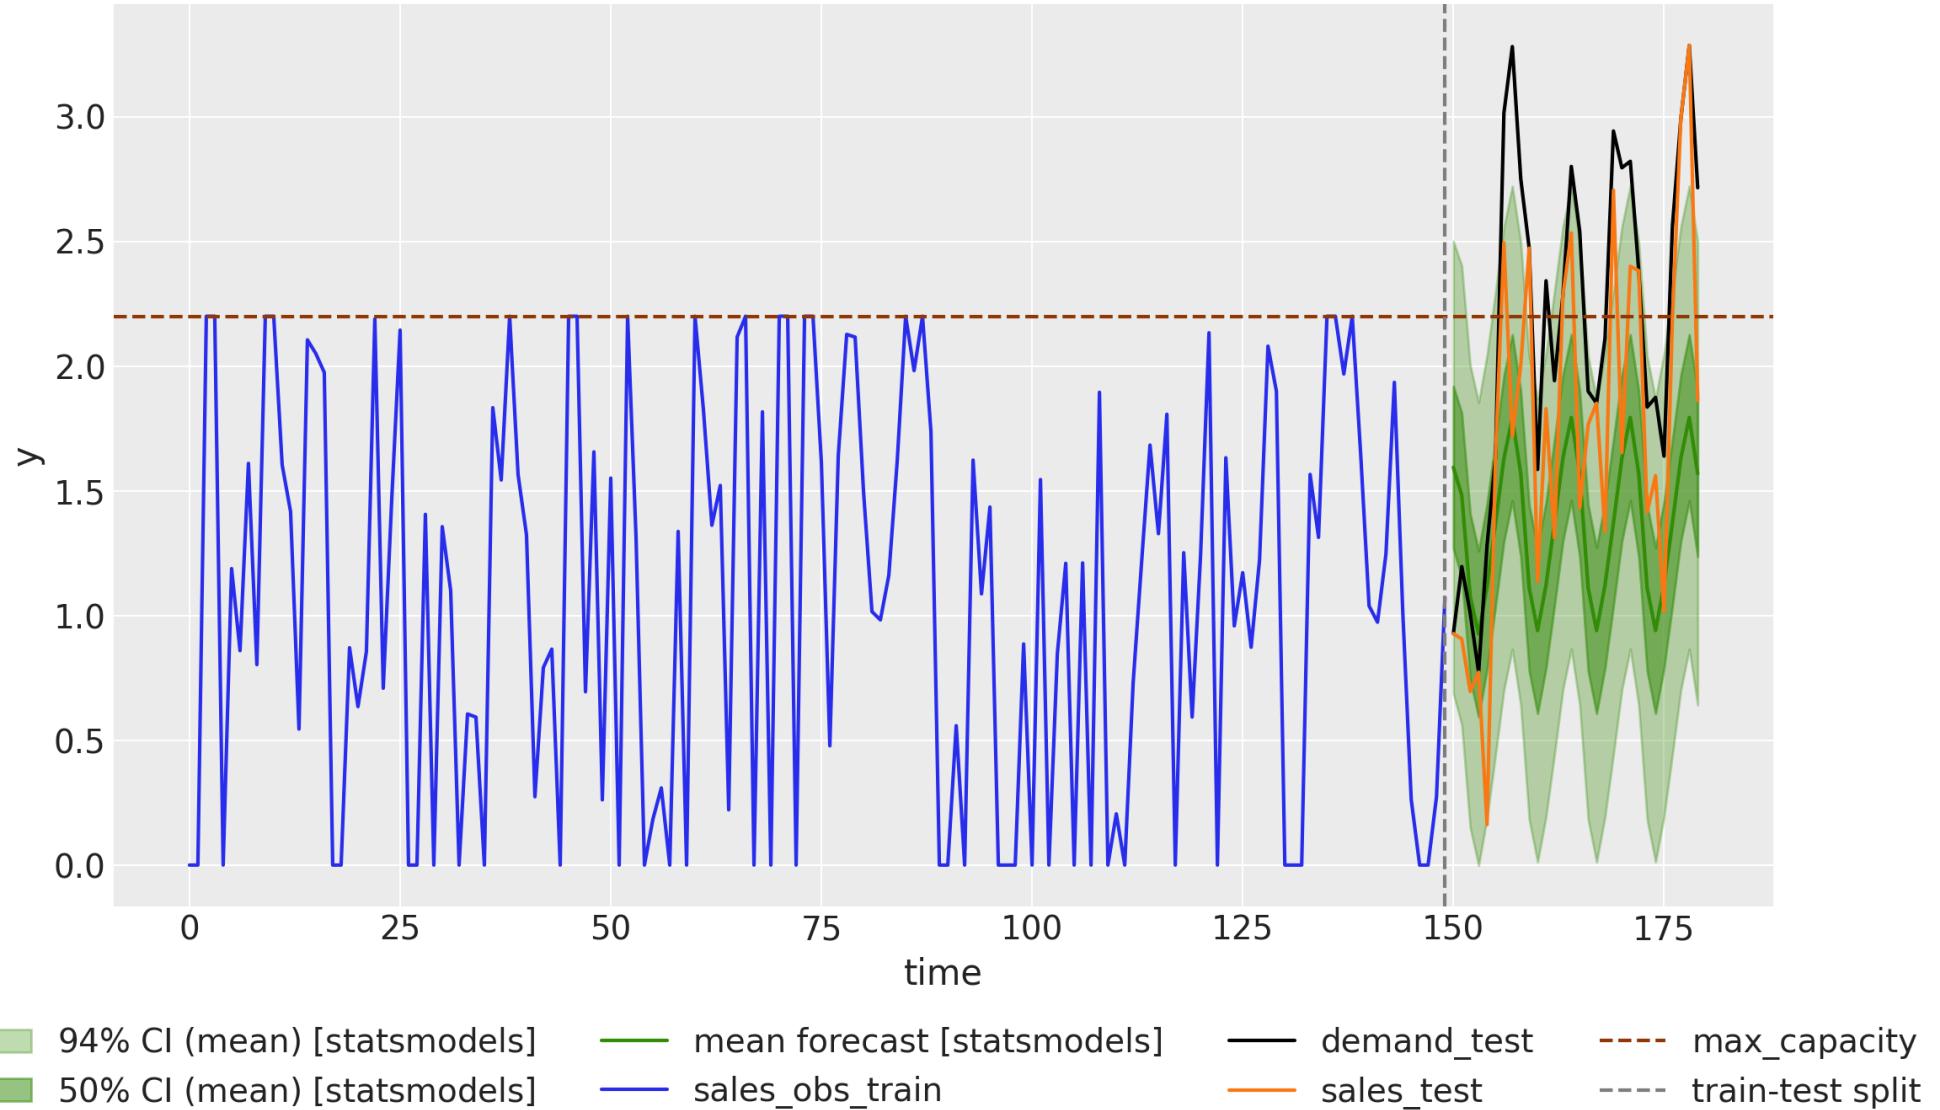
Forecasting Unseen Demand

Demand and Sales Simulation

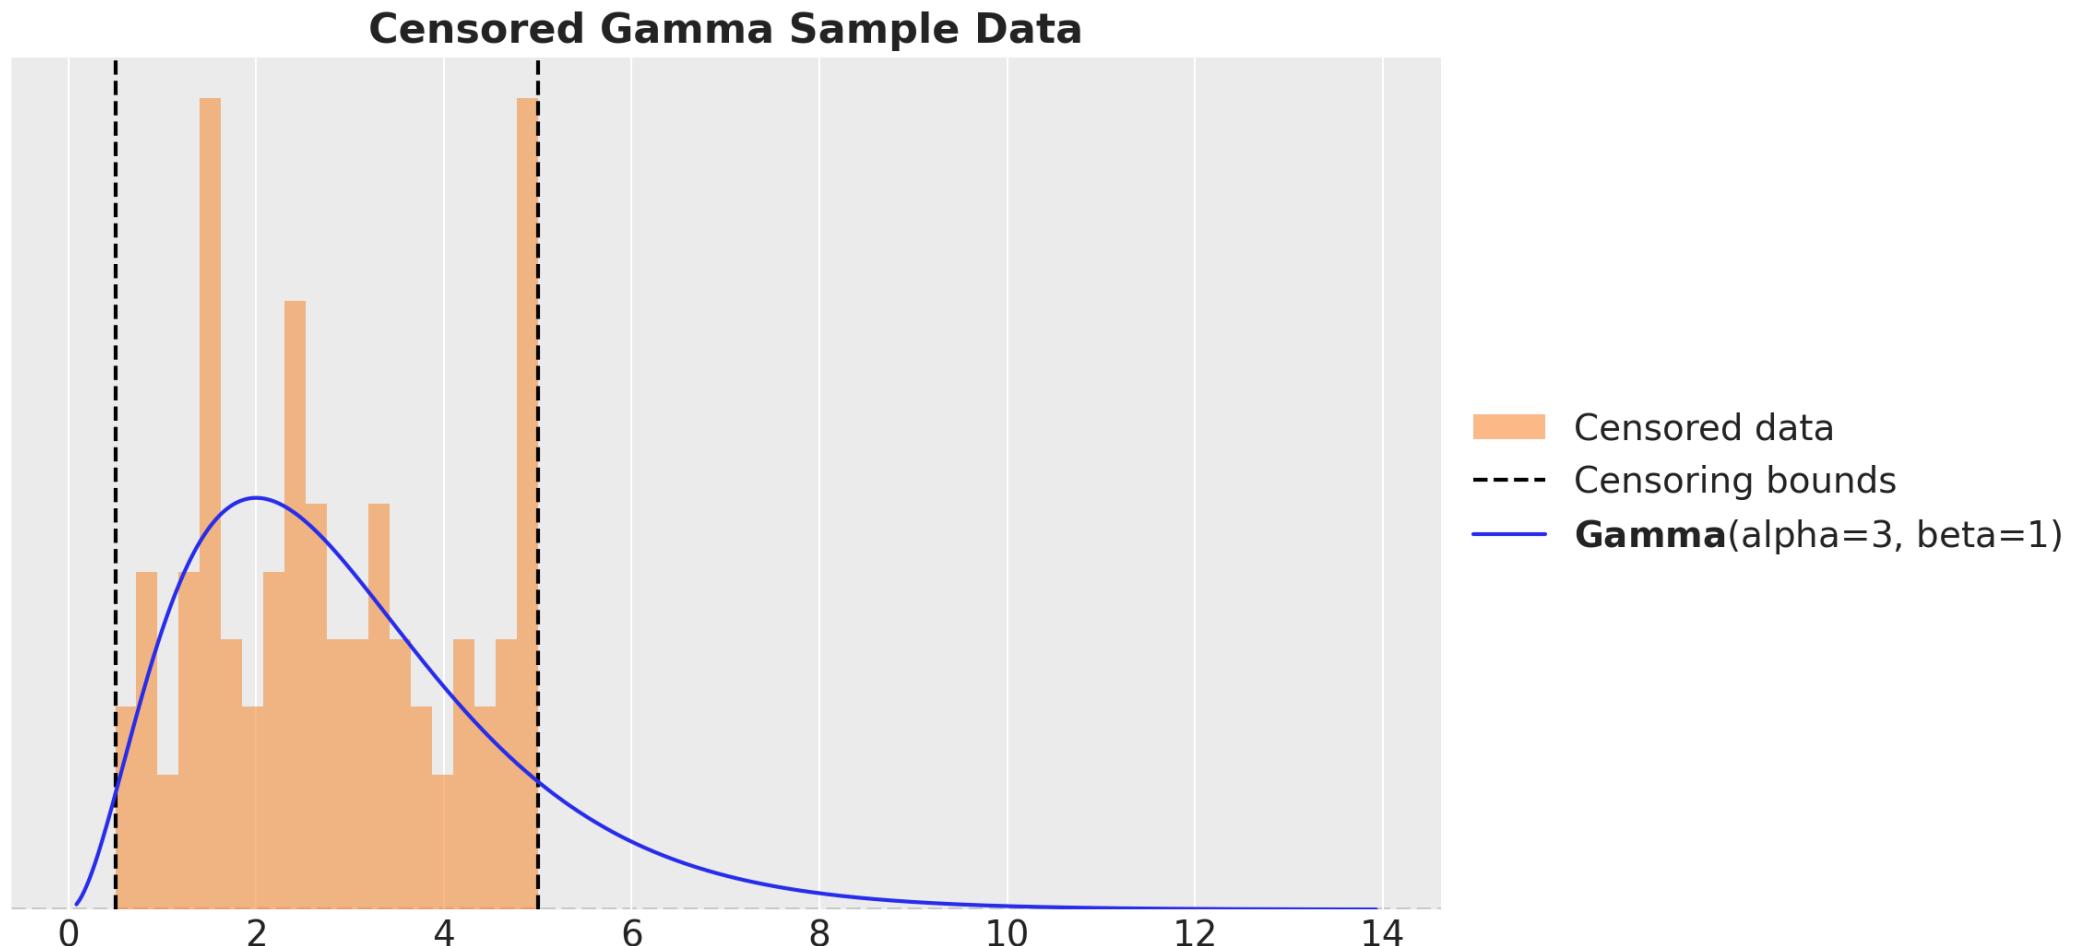


ARIMA Model

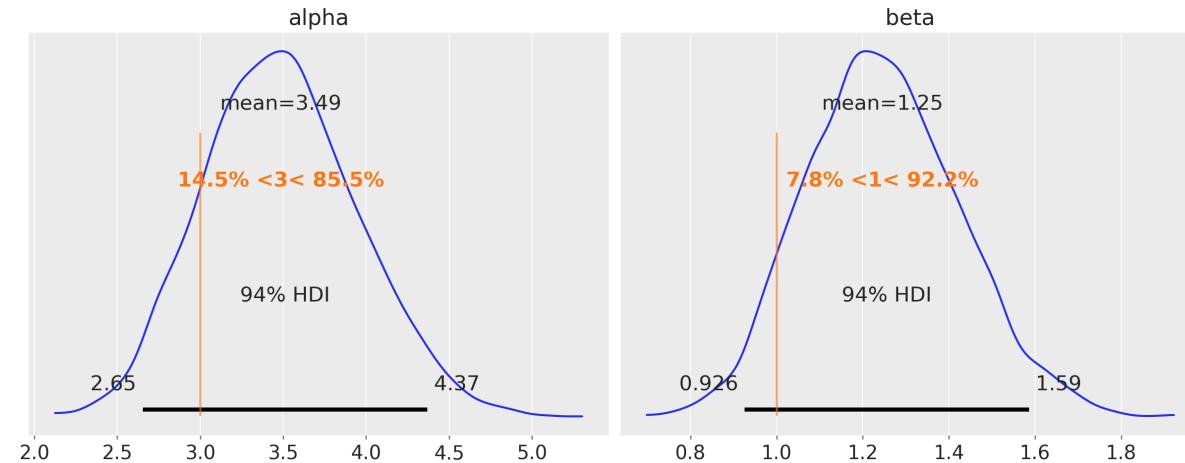
ARIMAX(2, 0, 0) + Fourier Seasonality Model Forecast (Statsmodels)



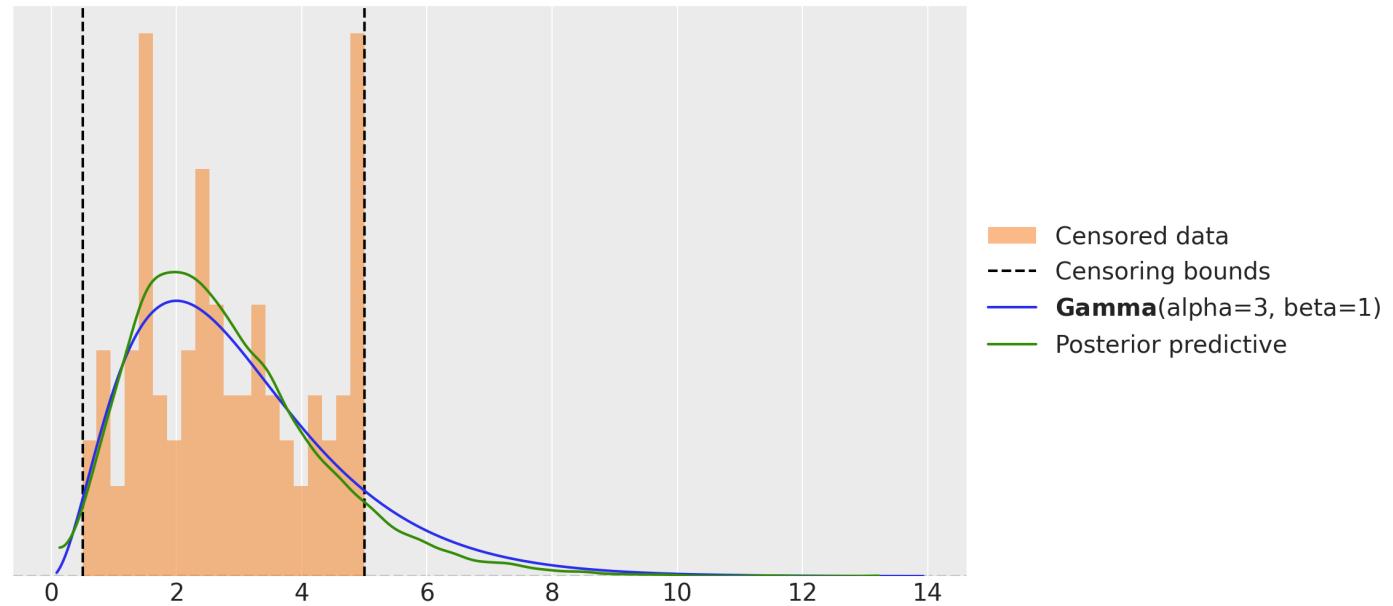
Censored Distributions

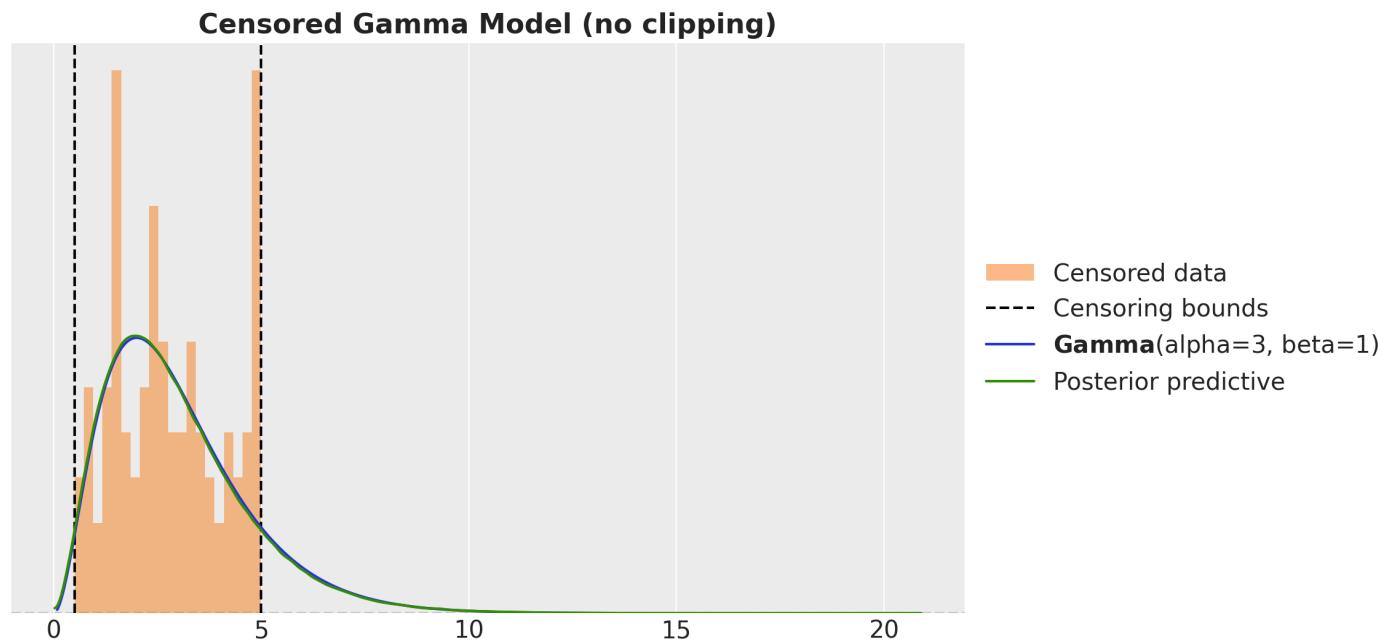
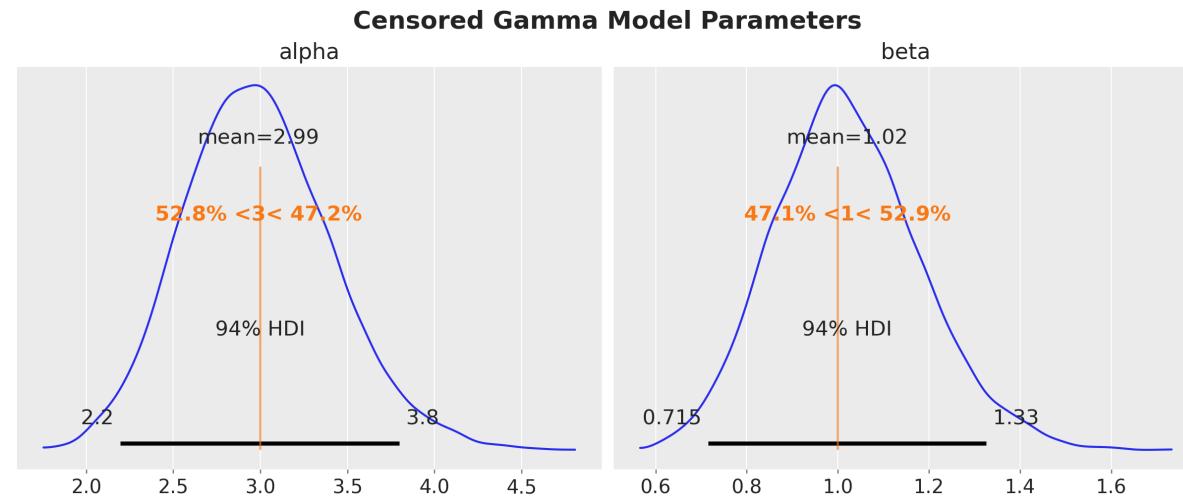


Gamma Model Parameters



Gamma Model





Censored Likelihood

```

1 def censored_normal(loc, scale, y, censored):
2     distribution = dist.Normal(loc=loc, scale=scale)
3     ccdf = 1 - distribution.cdf(y)
4     numpyro.sample(
5         "censored_label",
6         dist.Bernoulli(probs=ccdf).mask(censored == 1),
7         obs=censored
8     )
9     return numpyro.sample("pred", distribution.mask(censored != 1))

```

Change likelihood distribution in a time-series model:

```

1 ## Transition function for AR(2)
2 def transition_fn(carry, t):
3     y_prev_1, y_prev_2 = carry
4     ar_part = phi_1 * y_prev_1 + phi_2 * y_prev_2
5     pred_mean = mu + ar_part + seasonal[t]
6     # Censored likelihood
7     pred = censored_normal(pred_mean, sigma, y[t], censored[t])
8     return (pred, y_prev_1), pred

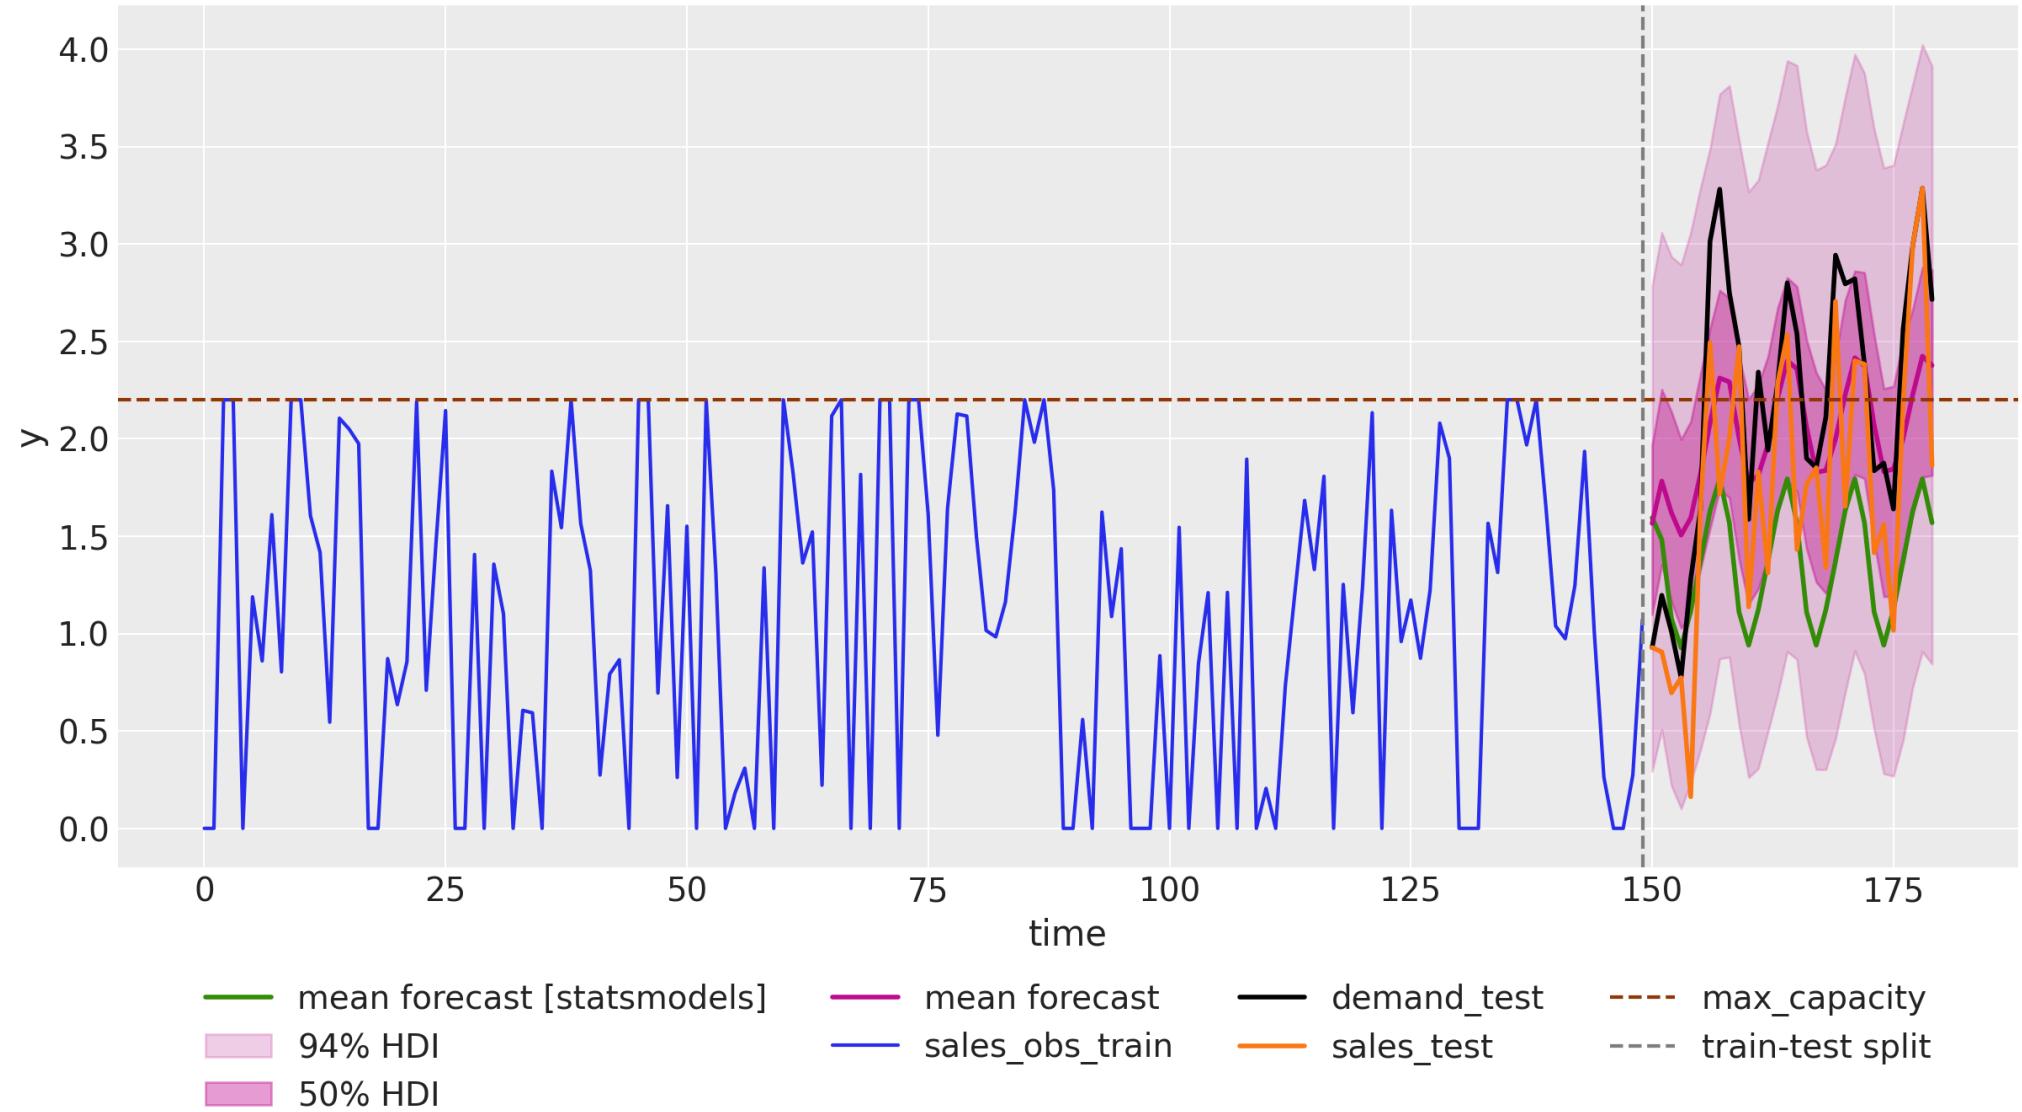
```



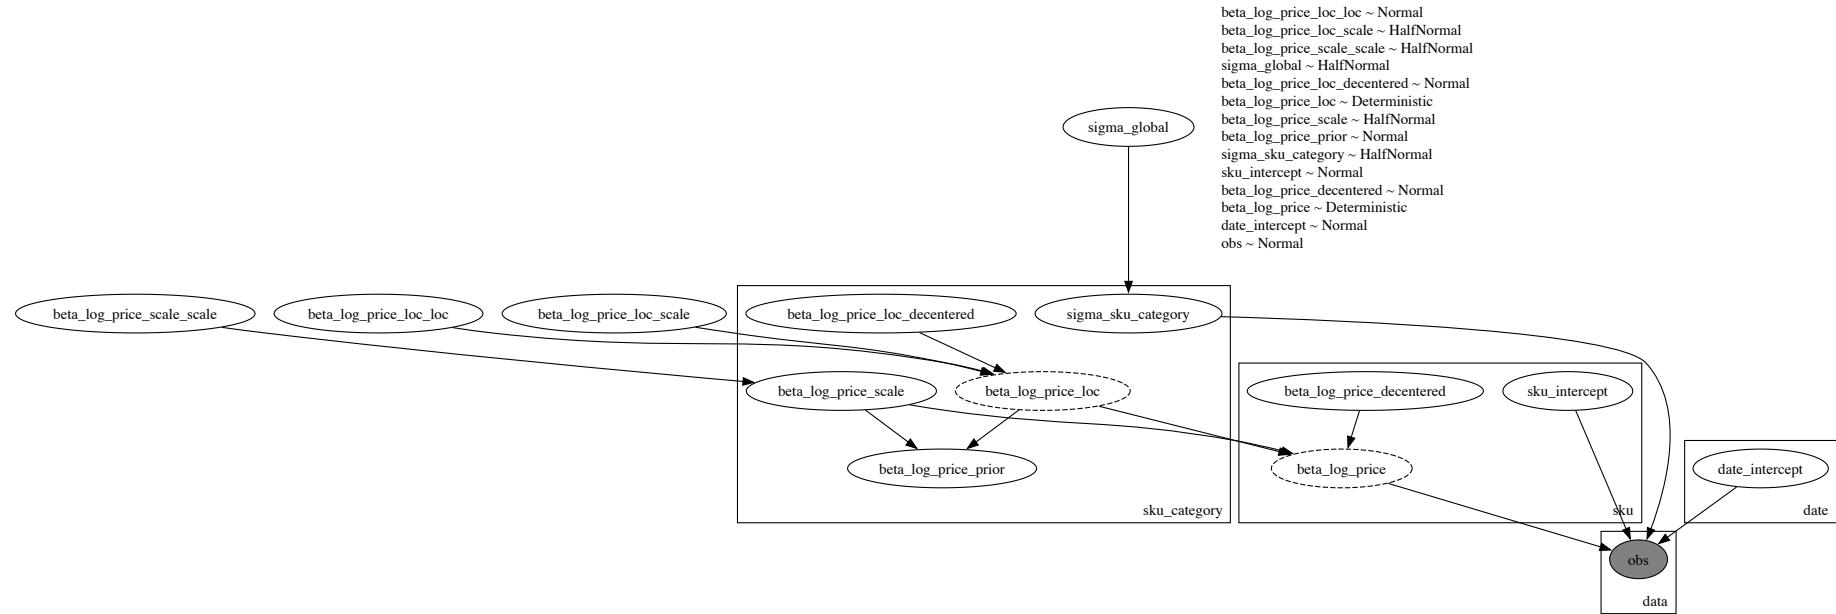
Censored Time Series Forecast



Censored AR(2) + Fourier Modes - Forecast



Hierarchical Pricing Elasticity Models

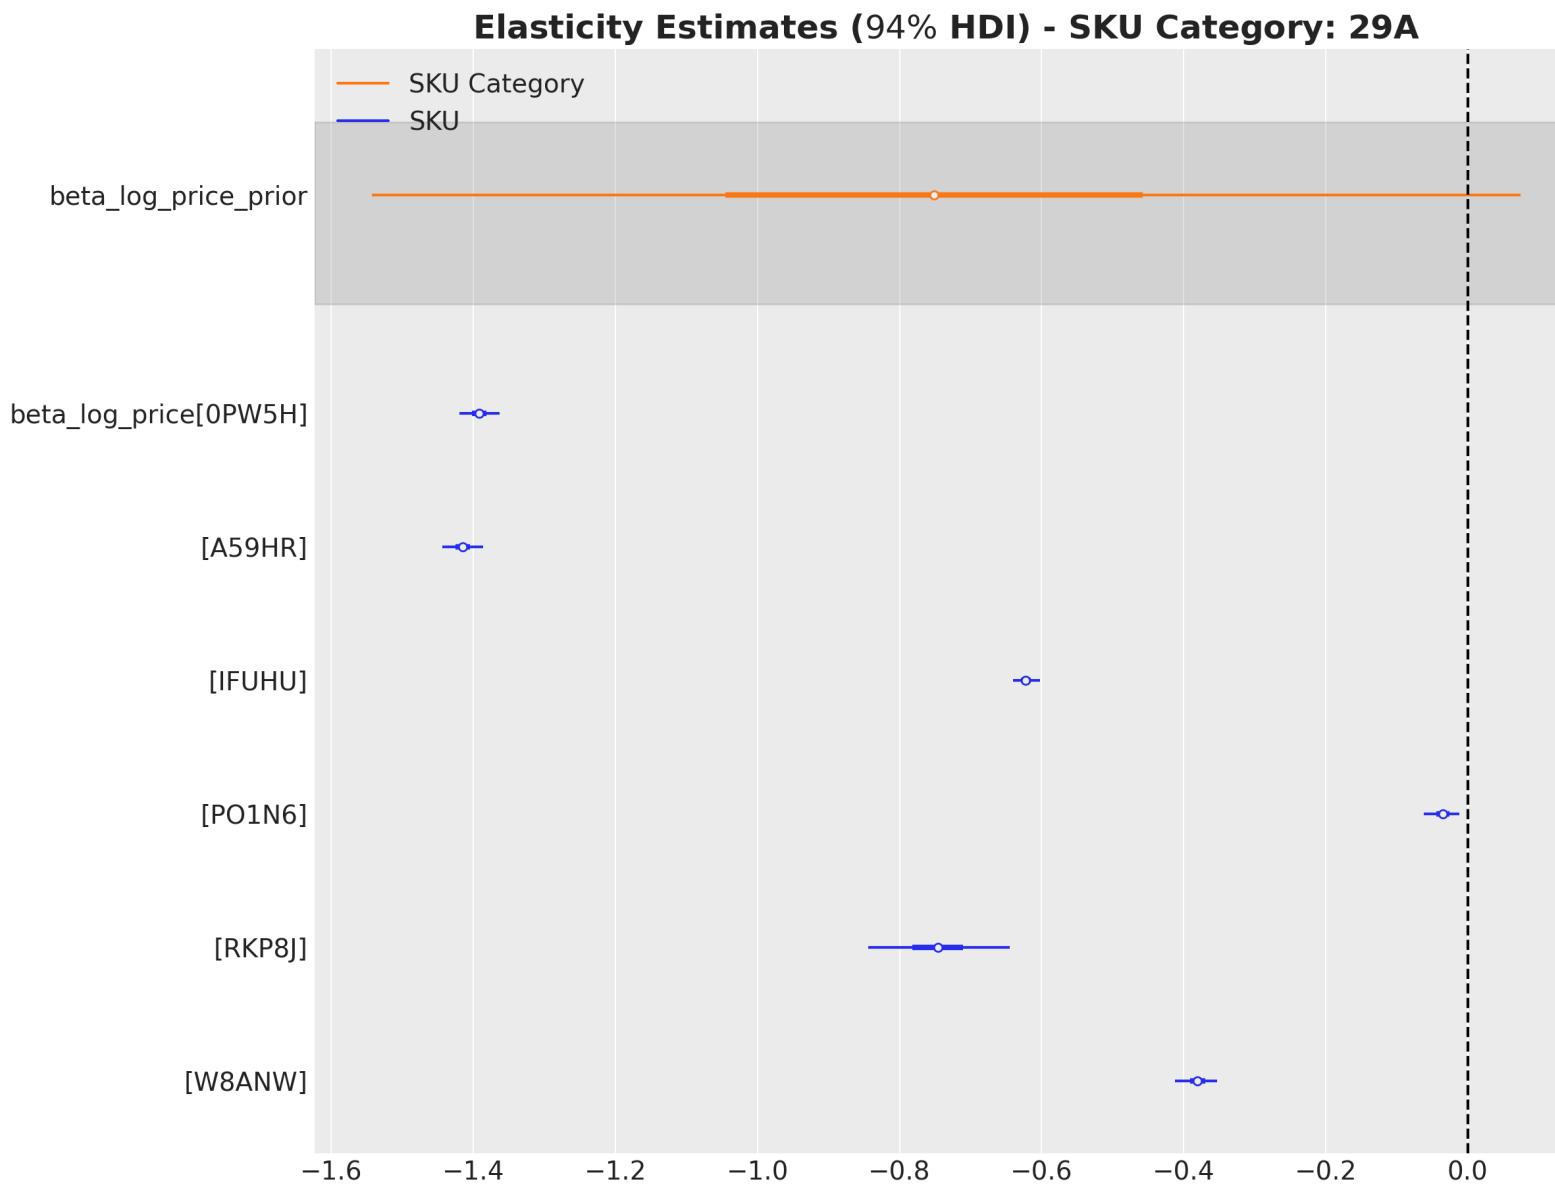


Idea 😎

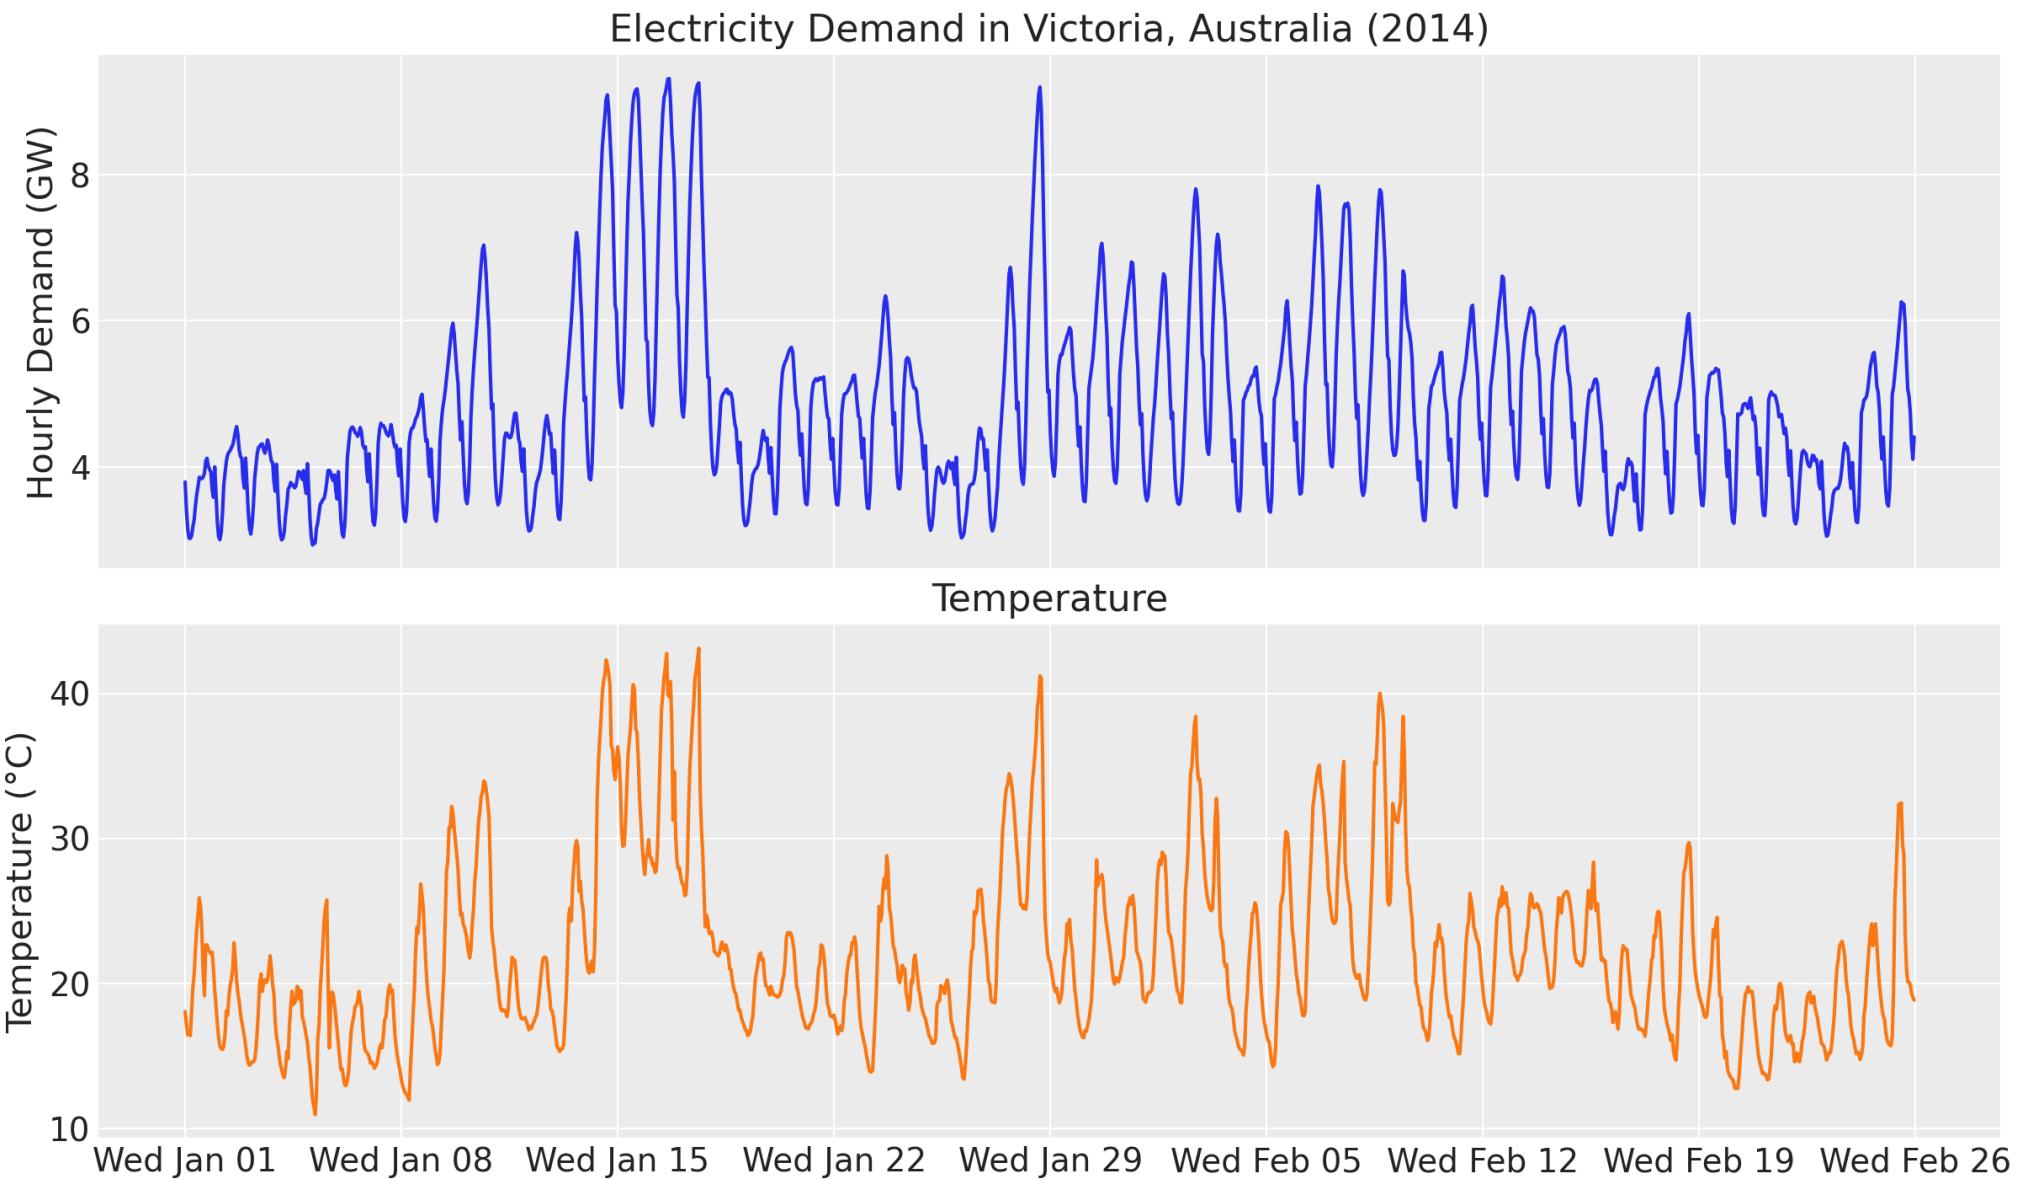
Use a hierarchical structure to regularize the demand elasticity parameters.



Hierarchical Pricing Elasticity Models

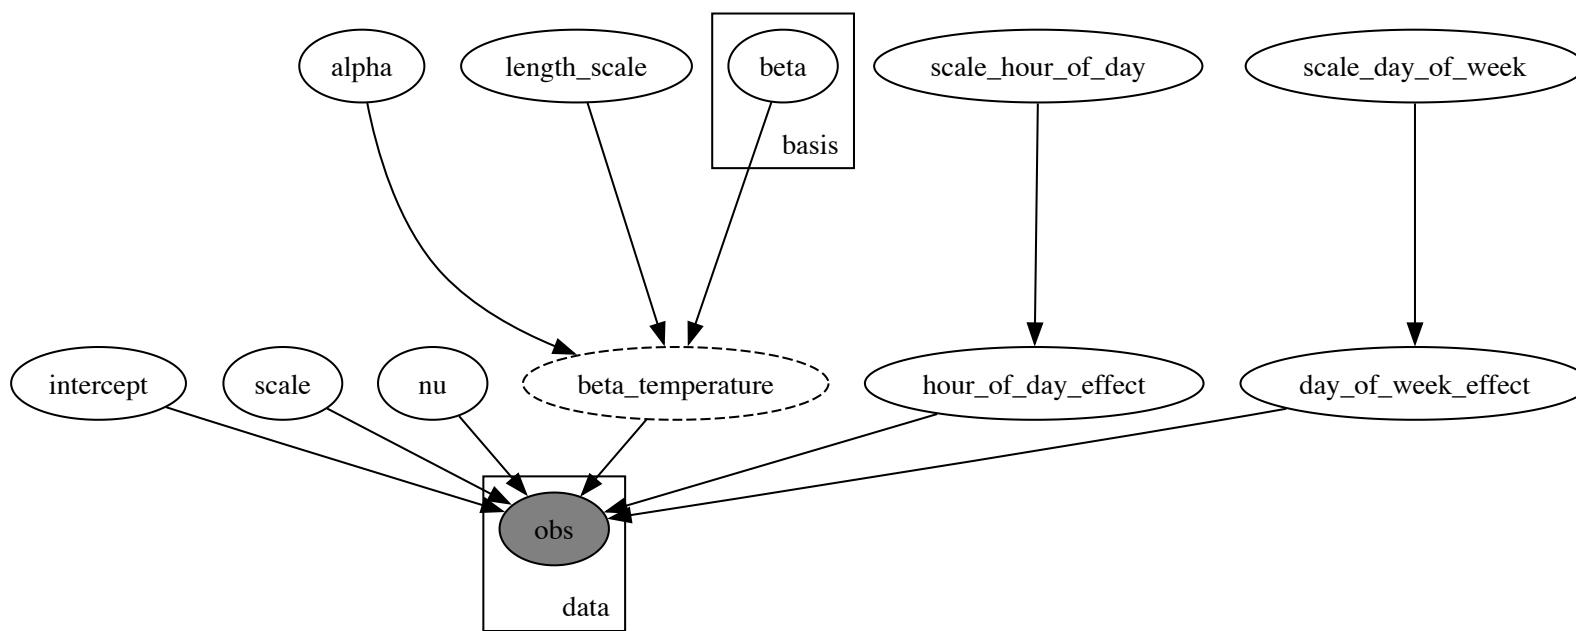


Dynamic Time-Series Model



Dynamic Coefficients

Hilbert Space Gaussian Processes for Dynamic Coefficients



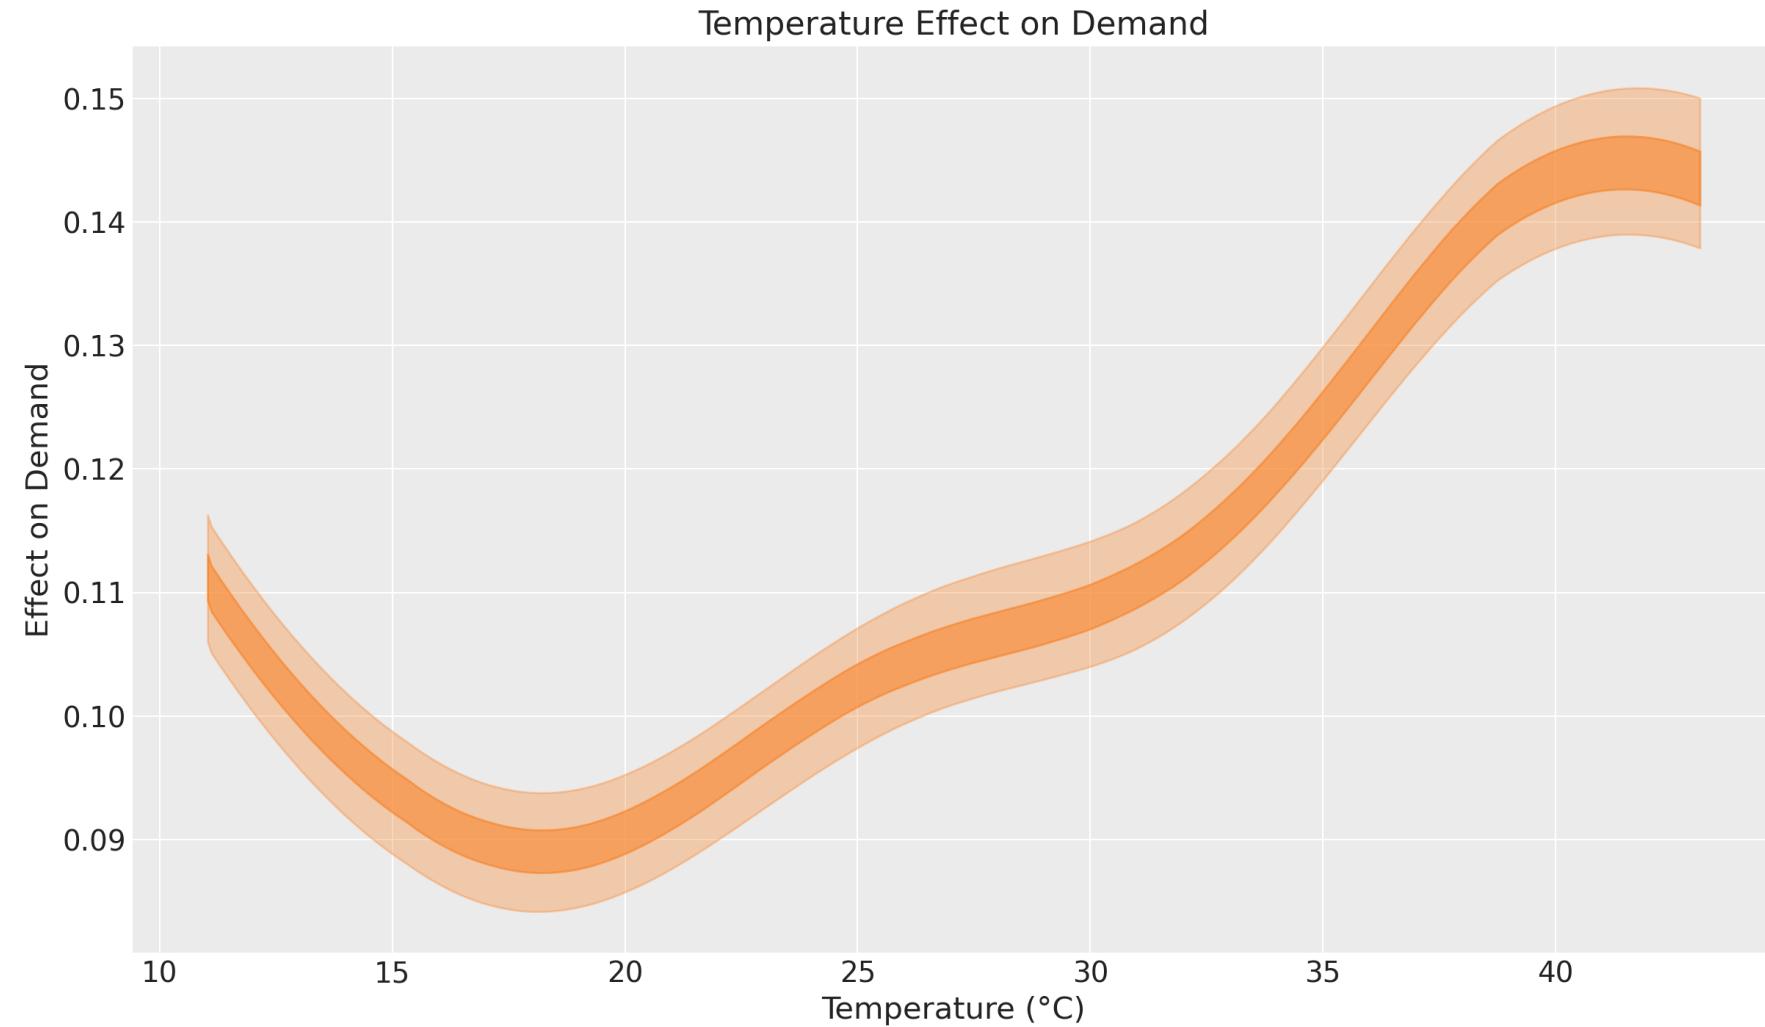
```

intercept ~ Normal
alpha ~ InverseGamma
length_scale ~ InverseGamma
scale ~ HalfNormal
nu ~ Gamma
beta ~ Normal
beta_temperature ~ Deterministic
scale_hour_of_day ~ HalfNormal
hour_of_day_effect ~ ZeroSumNormal
scale_day_of_week ~ HalfNormal
day_of_week_effect ~ ZeroSumNormal
obs ~ StudentT

```



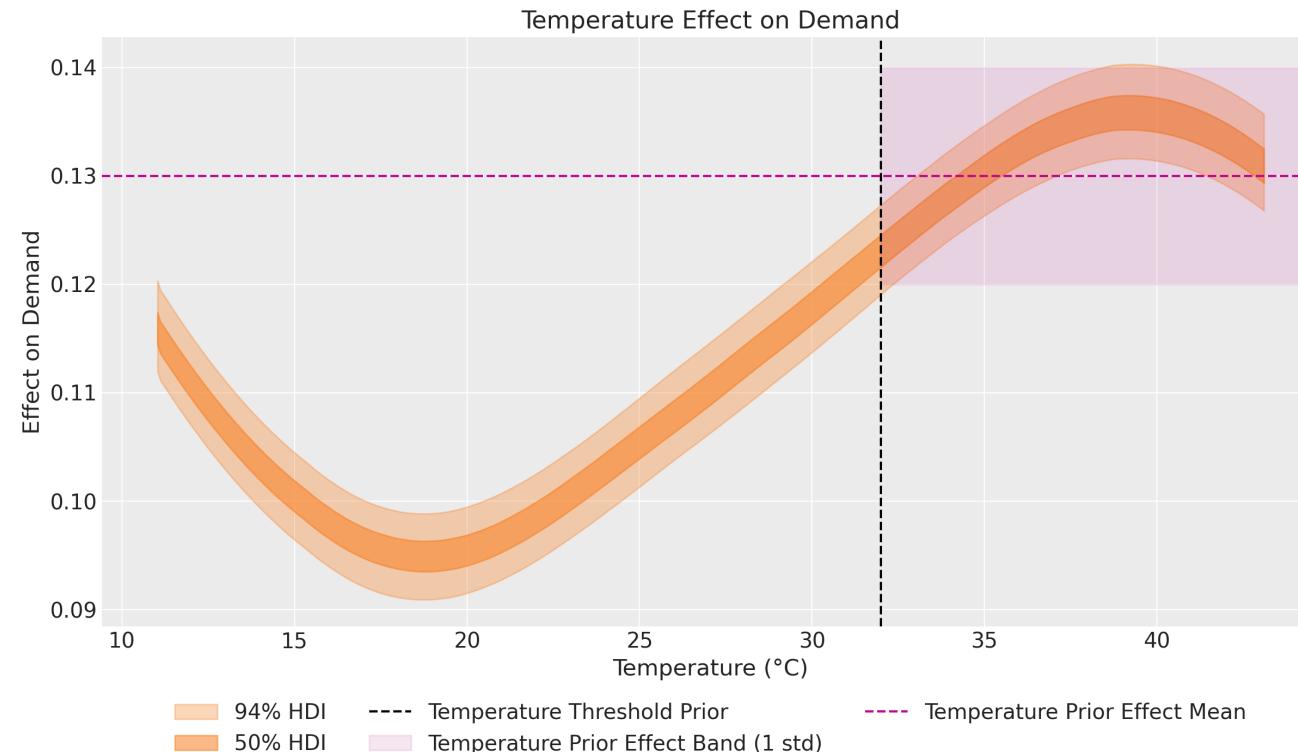
Inferring Effect of Temperature on Electricity Demand



Calibrating a Demand Model

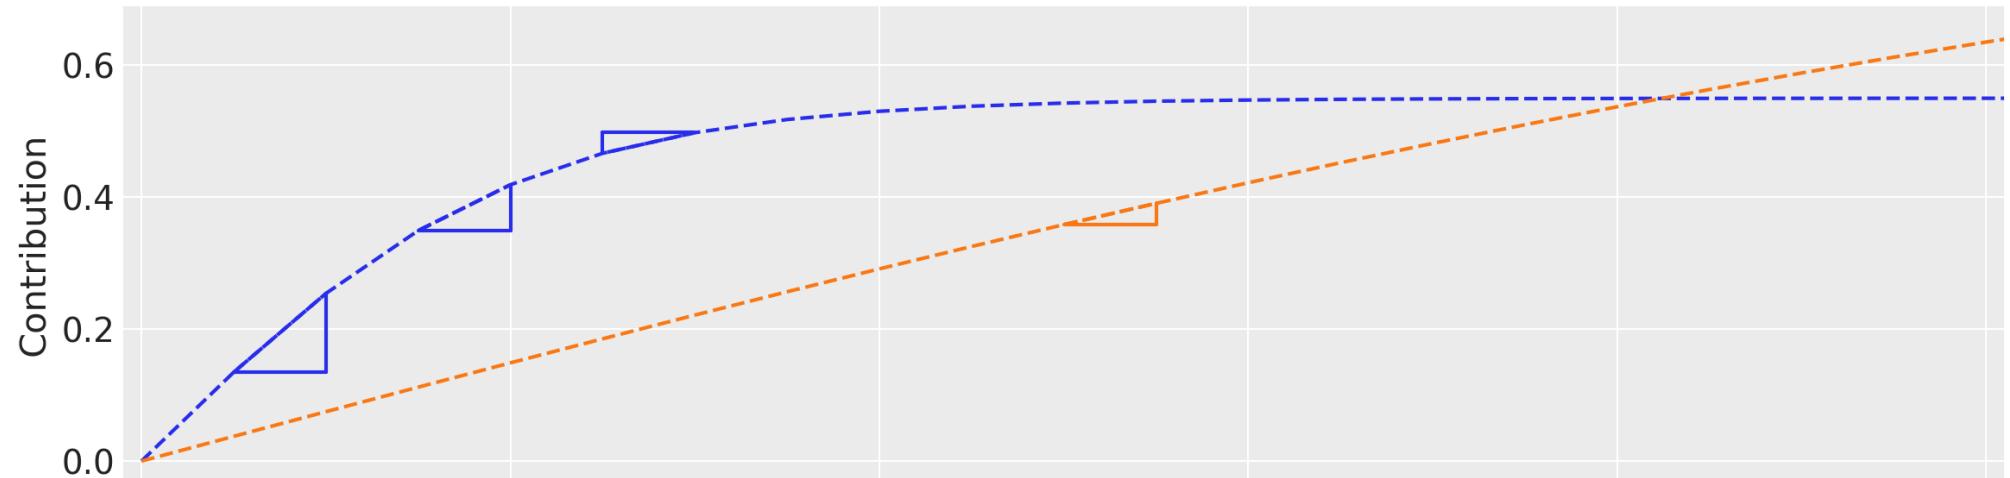


Let us assume that we know from domain knowledge that the effect of temperature on demand over 32°C is somehow stable at around a value of 0.13.

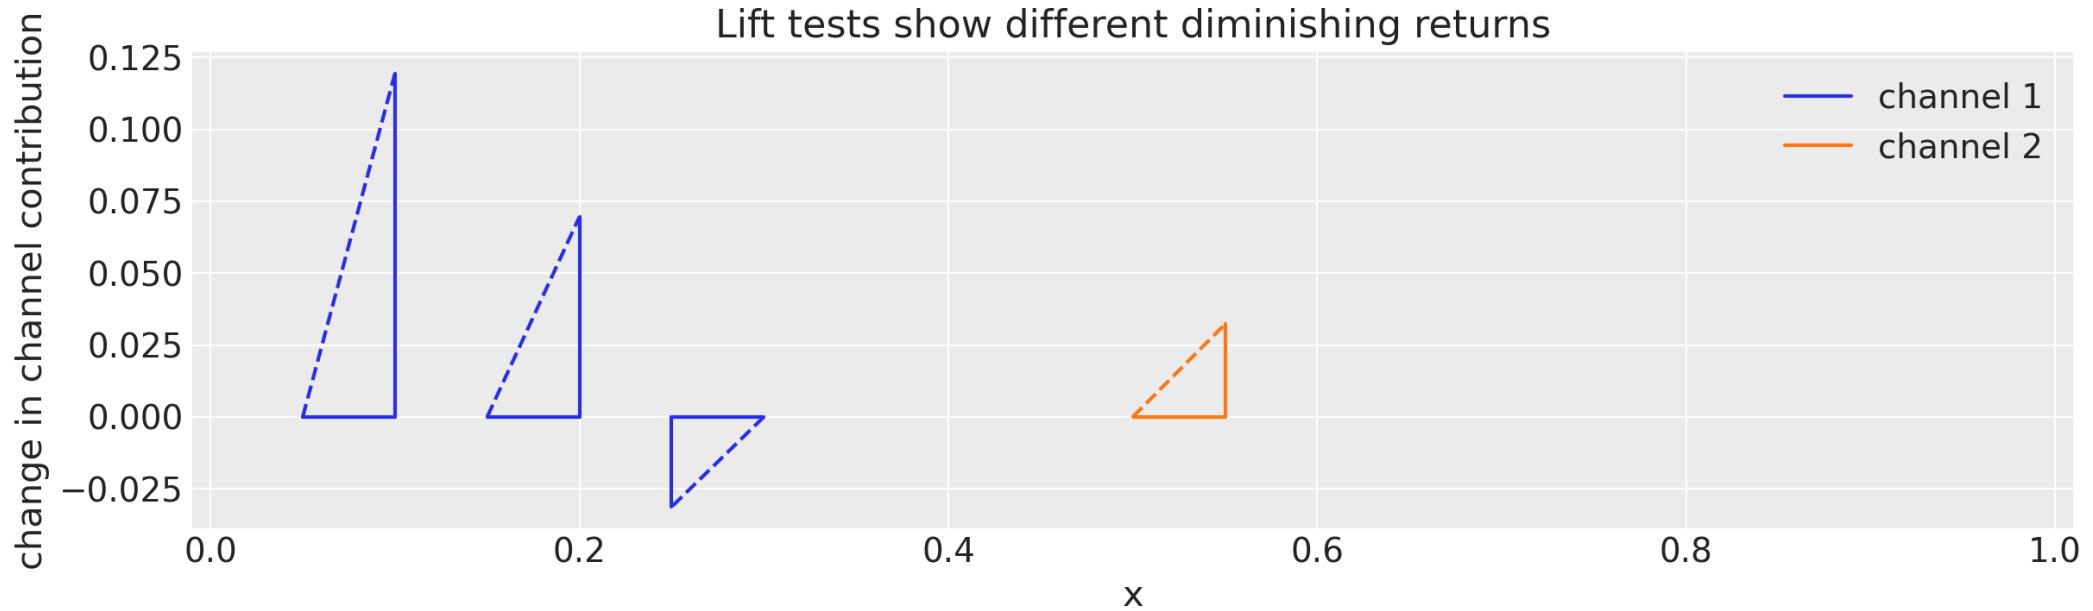


MMM Calibration with Lift Tests

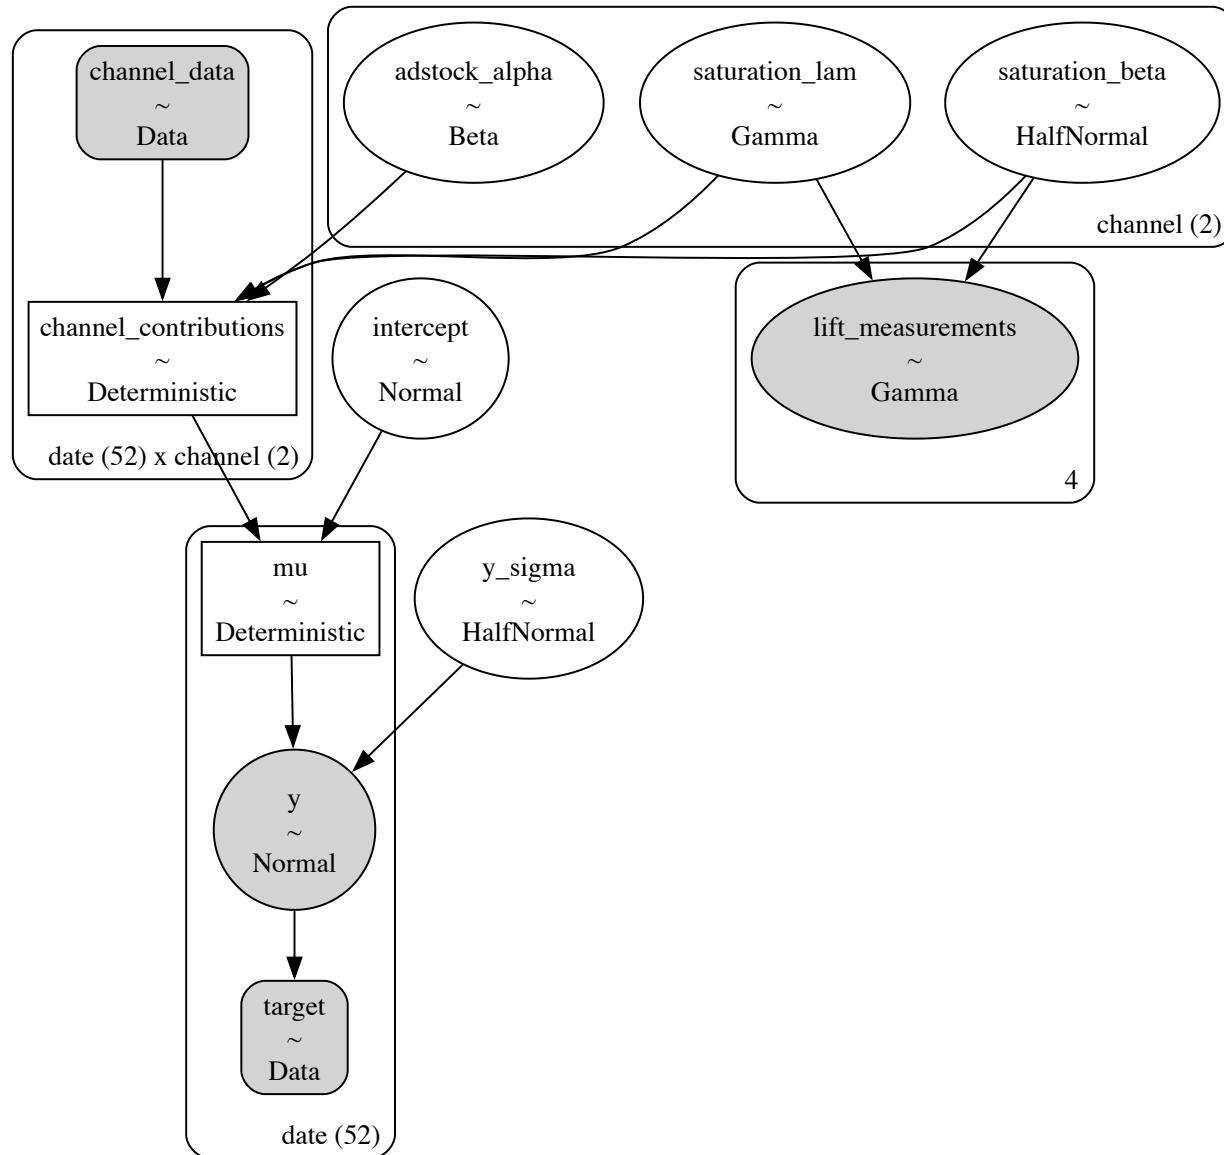
Lift tests results shown on top of actual curves



Lift tests show different diminishing returns



MMM Calibration with Lift Tests



References



NumPyro Examples

- Introduction to Stochastic Variational Inference with NumPyro ★
- Bayesian Censoring Data Modeling
- Croston's Method for Intermittent Time Series Forecasting in NumPyro
- Demand Forecasting with Censored Likelihood
- Electricity Demand Forecast: Dynamic Time-Series Model with Prior Calibration
- From Pyro to NumPyro: Forecasting Hierarchical Models - Part I
- From Pyro to NumPyro: Forecasting Hierarchical Models - Part II
- Hacking the TSB Model for Intermediate Time Series to Accommodate for Availability Constraints
- Hierarchical Exponential Smoothing Model
- Hierarchical Pricing Elasticity Models
- Notes on Exponential Smoothing with NumPyro



References



Packages:

- [Nixtla](#)
- [Prophetverse](#)
- [Pyro - Forecasting](#)
- [Statsmodels - Time Series Models](#)
- [TimeSeers](#)
- [GluonTS](#)
- [Zalando: PyTorchTS](#)



References



Other Blogposts

- Finally! Bayesian Hierarchical Modelling at Scale
- Modeling Anything With First Principles: Demand under extreme stockouts
- PyMC-Marketing: Lift Test Calibration
- PyMC Labs: Unobserved Confounders, ROAS and Lift Tests in Media Mix Models
- PyMC Labs: Probabilistic Time Series Analysis - Opportunities and Applications

Books

- Forecasting: Principles and Practice, the Pythonic Way
- Forecasting and Analytics with the Augmented Dynamic Adaptive Model (ADAM)

Papers



Thank you!

juanitorduz.github.io

